

1. What is Kafka?

Apache Kafka is a open source distributed event streaming platform. Lets divide these terminologies,

Herer streaming is of two types.

Creating Real Time Stream: Sending the stream of continuous data from Paytm to Kafka server is known as Creating Real Time Streaming

For example I am booking the ticket from Paytm. Here I am not only the user booking the ticket from world wise Paytm may receive multiple requests for different users to book the tickets on each minute or second.

Processing Real Time Stream: The process of continuously listening and processing the data from Kafka server is known as processing real time stream.

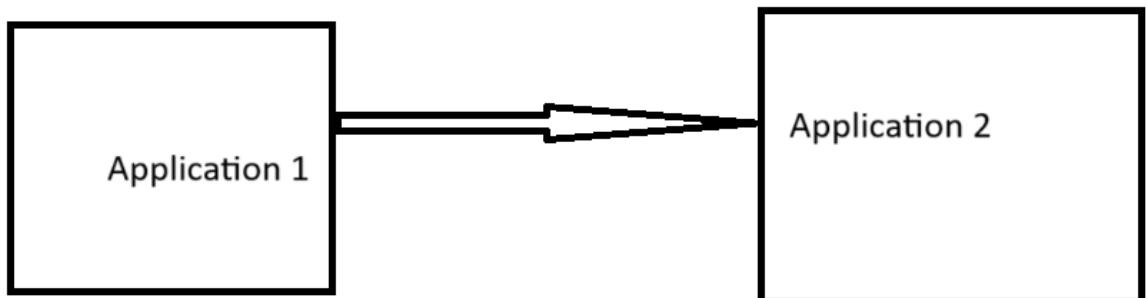
Distributed: In micro services world everything is load is distributed across the regions to balance the load to avoid the downtime. Likewise Kafka is also distribute Kafka servers in multiple regions to avoid the downtime and load balance. If any of the Kafka server is down another will come into picture and handle the downtime.

2. Where does Kafka come from?

Kafka was originally developed by LinkedIn and it was subsequently open sourced in early 2011 to Apache.

3. Where do we need Kafka?

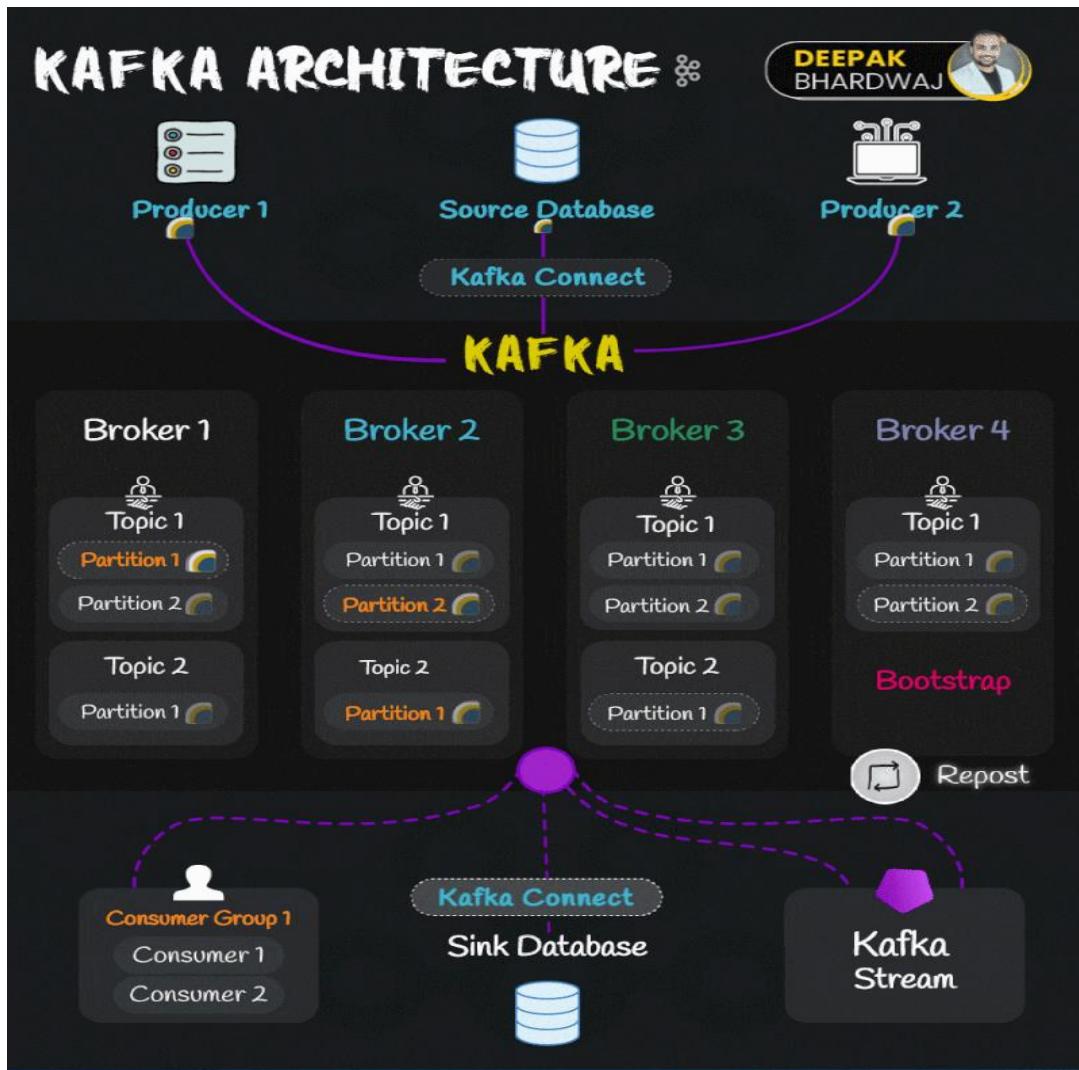
Consider an example I went to vacation and I got an important parcel from bank. As I was not available at home parcel may went back to provider he tried 3 more attempts to deliver the parcel and I may miss the parcel. In this case if I put the parcel box at my gate postman will put the parcel in and I wont miss the parcel. Here post box will work as middle man between me and postman. Hence post box will be the Kafka.



In real-time example I am sending the data from application 1 to application 2. In case application 2 is down and not able to receive the data from application 1. Hence data will be

lost from Application 1 to Application 2 like postman missed to deliver the letter to me. As I added the post box to receive the letters at my home in case of my unavailability here also I need some middle man to store the data in case of applications is down. Here the middle man is Kafka.

4. Explain Kafka Internal workflow?

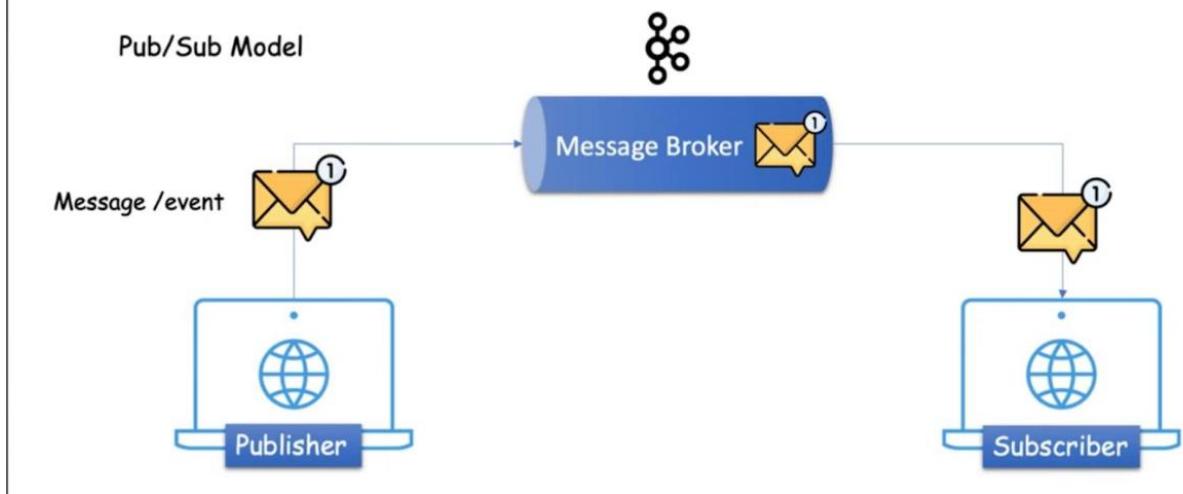


- ↳ Producers send data to specific Topics in the Kafka Cluster.
- ↳ Data in Topics is divided into Partitions for load balancing.
- ↳ Leader Partitions manage data integrity, while Follower Partitions provide backup.
- ↳ Consumers pull data, enabling real-time analytics and decision-making.

5. How does Kafka work?

Kafka works on pub sub model

How does it work (High-level overview)

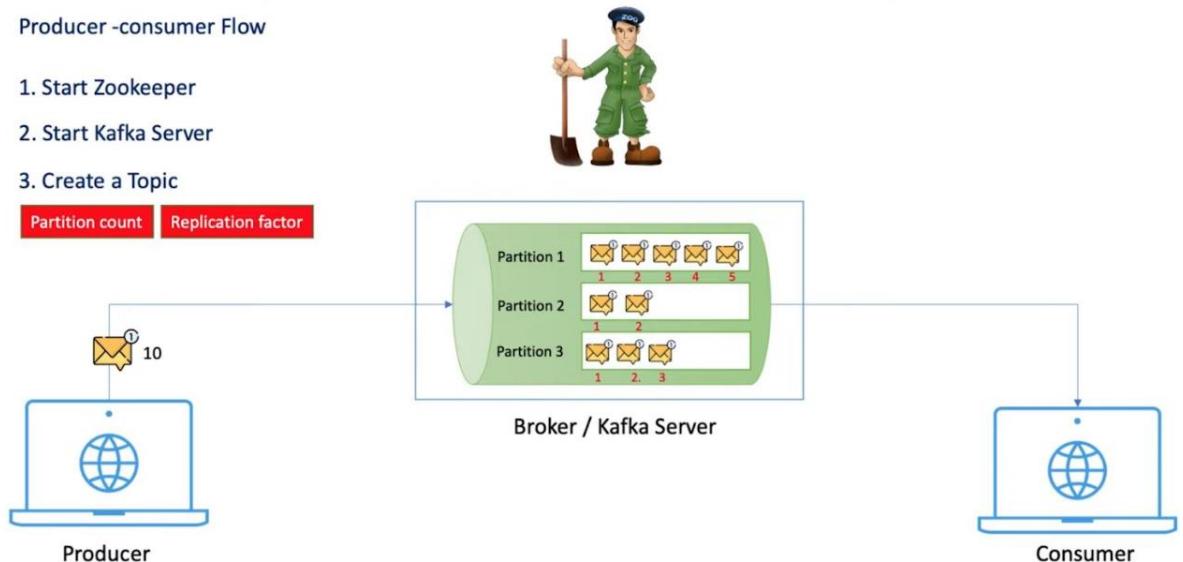


Publisher: He is the person who will send the messages.

Message Broker: A broker is nothing but a just server. In simple world a broker is just an intermediate entity that helps in message exchanges between producer and consumer.

Subscriber. Who will consume the messages from Kafka.

Message Broker Internal View



6. Explain Kafka Architecture?

Topics: A topic is a category or feed name to which records are sent. Topics are split into partitions for scalability and fault tolerance.

Producers: Producers are clients that publish (write) records to Kafka topics. Producers can choose which partition to write to based on key values or round-robin distribution.

Consumers: Consumers are clients that read records from Kafka topics. Each consumer belongs to a consumer group, and each record is delivered to only one consumer in a group.

Brokers: Brokers are Kafka servers that store data and serve client requests. A Kafka cluster is made up of multiple brokers to ensure fault tolerance and scalability.

Partitions: Each topic is divided into partitions, which are ordered, immutable sequences of records. Partitions allow Kafka to parallelize data consumption and storage.

Replicas: Kafka replicates partitions across multiple brokers for fault tolerance. Each partition has one leader and multiple followers. The leader handles all reads and writes for the partition, while followers replicate the data.

Zookeeper: Zookeeper is a centralized service for maintaining configuration information, naming, and distributed synchronization. It is used by Kafka to manage cluster metadata and broker coordination.

How Kafka Works

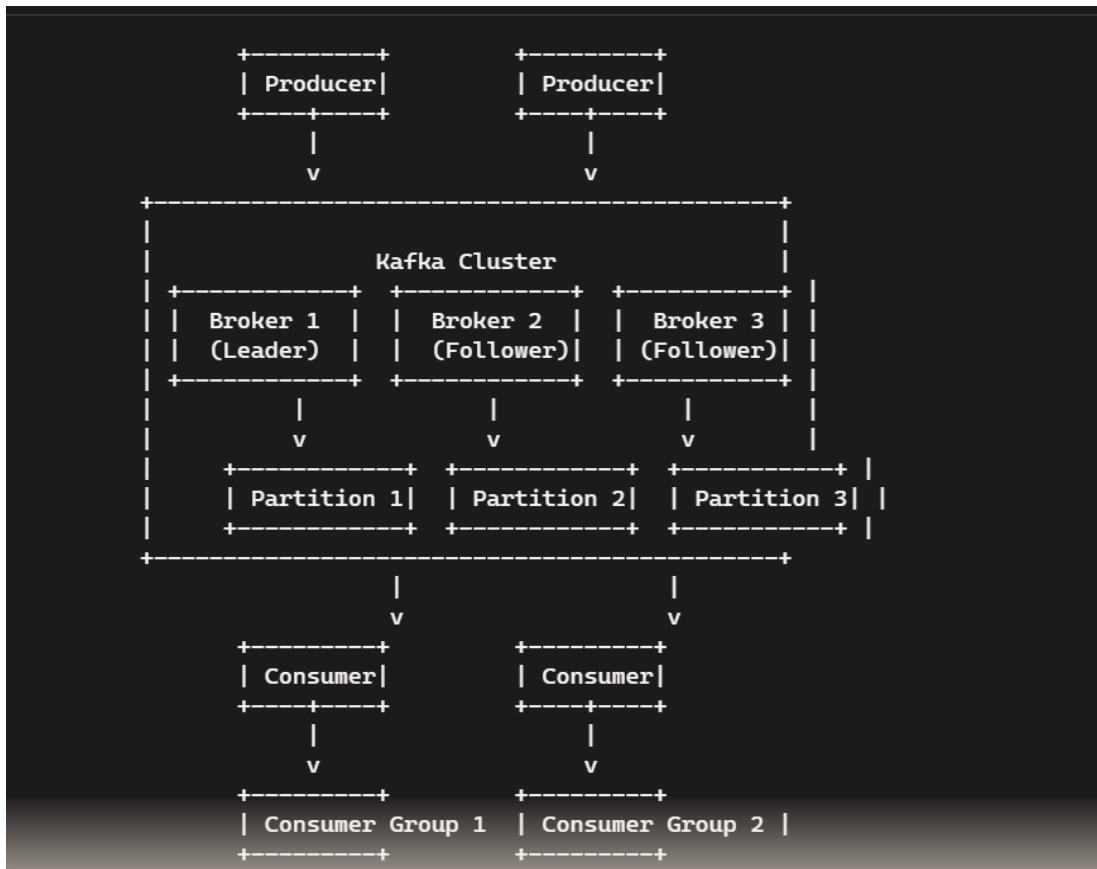
Producers send records to Kafka topics: Producers write records to specific topics. Each record consists of a key, value, timestamp, and metadata.

Partitions and replication: Each topic is divided into partitions, and each partition is replicated across multiple brokers. The leader broker for each partition handles read and write requests, while follower brokers replicate the data.

Consumers read records from Kafka topics: Consumers read records from topics by subscribing to one or more topics. Consumers in the same consumer group share the load of reading records from partitions.

Offset management: Kafka maintains the offset, which is the position of a consumer in a partition. Consumers use offsets to track their progress and ensure they don't miss or duplicate records.

Fault tolerance and scalability: Kafka ensures fault tolerance through partition replication. If a leader broker fails, a follower takes over as the leader. Kafka scales horizontally by adding more brokers and partitions to handle increased load and storage requirements.



In this architecture:

- Producers send data to Kafka topics, which are divided into partitions.
- Brokers manage these partitions, with one broker acting as the leader for each partition.
- Consumers read data from these partitions, coordinated through consumer groups.

7. How to start the Kafka in local system?

- ✓ Start Zookeeper
- ✓ Start Kafka Server
- ✓ Start Producer
- ✓ Start Consumer

Zookeeper default port: **2181**

Kafka Server/Broker default port: **9092**

8. Can we run Kafka with out zookeeper?

Apache Kafka has officially deprecated **ZooKeeper** in version 3.5. The introduction of the KRaft mode feature not only enhances the scalability capabilities of Kafka but also simplifies the learning curve by eliminating the need to manage ZooKeeper.

9. What will happen if zookeeper fails in Kafka architecture?

If Zookeeper fails in an Apache Kafka cluster, the impact depends on whether the failure is temporary (e.g., one Zookeeper node goes down) or complete (entire Zookeeper ensemble fails). Here's what happens in both scenarios:

Single Zookeeper Node Failure (in a multi-node ensemble)

If you have a multi-node Zookeeper ensemble (typically 3, 5, or more nodes), Kafka can continue working as long as a majority (quorum) of Zookeeper nodes remain operational. The remaining healthy Zookeeper nodes continue to manage leader election, broker metadata, and topic partition state.

Complete Zookeeper Failure

If the entire Zookeeper ensemble fails (all nodes go down), Kafka brokers continue running, but the cluster starts experiencing issues:

New broker nodes cannot register or join the cluster.

Leader elections for partitions stop—if a broker (leader) crashes, no new leader can be elected.

No administrative operations (like topic creation/deletion) can be performed.

Producers and consumers might still function as long as the current partition leaders remain up, but if a leader broker crashes, that partition becomes unavailable.

Recovery from Zookeeper Failure

If Zookeeper nodes restart and regain quorum, Kafka resumes normal operations.

If Zookeeper data is lost or corrupted, you may need to restore it from a backup.

How to Prevent Kafka Disruptions Due to Zookeeper Failure?

Run a multi-node Zookeeper ensemble for high availability.

Monitor Zookeeper health using tools like Prometheus, Grafana, or Zookeeper's built-in status commands.

Migrate to KRaft (Kafka Raft)—newer versions of Kafka (from 2.8 onwards) introduce KRaft mode, which eliminates the need for Zookeeper by managing metadata within Kafka itself.

10. Explain about Kraft architecture?

Key Components in KRaft

KRaft Controllers: KRaft introduces controllers that manage all the metadata and cluster state. Unlike the ZooKeeper-based approach, these controllers use the Raft consensus algorithm for coordination and leader election.

Raft Consensus Protocol: This protocol ensures that all metadata updates are replicated across multiple controller nodes to provide fault tolerance. It helps in achieving consensus even when some nodes fail.

Brokers: Kafka brokers in KRaft mode no longer require ZooKeeper. They directly communicate with the KRaft controllers for metadata updates and cluster management.

KRaft Architecture Flow

Cluster Initialization: The cluster is initialized with one or more KRaft controllers. These controllers are responsible for managing the metadata and cluster state.

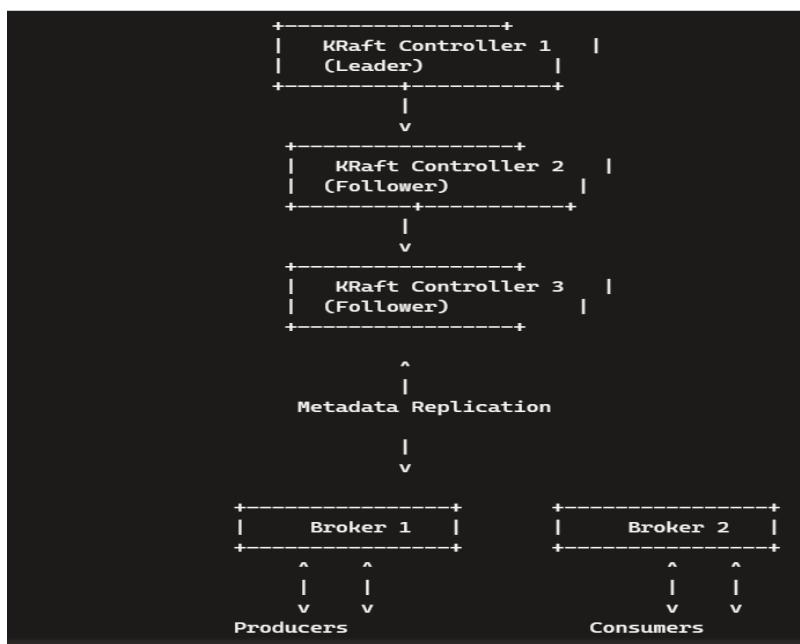
Leader Election: KRaft controllers use the Raft consensus protocol to elect a leader. The leader controller is responsible for making decisions and updating metadata.

Metadata Management: All metadata changes, such as topic creation, partition assignment, and broker registrations, are managed by the KRaft controllers. The leader controller updates the metadata and replicates it to follower controllers.

Broker Operations: Kafka brokers interact directly with the KRaft controllers to fetch metadata and update their state. They no longer need to communicate with ZooKeeper.

Client Interactions: Producers and consumers continue to interact with brokers as usual. They are unaware of the underlying changes in metadata management.

Fault Tolerance: If the leader controller fails, the Raft protocol ensures that a new leader is elected from the remaining controllers, maintaining high availability.



11. How the leader elections happens in Kafka?

Every broker has information about the list of topics (and partitions) and their leaders which will be kept up to date by the zoo keeper whenever the new leader is elected or when the number of partition changes.

Thus, when the producer makes a call to one of the brokers, it responds with this information list. Once the producer receives this information, it caches this and uses it to connect to the leader. So next time when it wants to send the message to that particular topic (and partition) it will use this cached information.

Lets assume there was only one leader and there are no replicas for that topic/partition duo and it got crushed. In this case it will try to connect to that leader and it fails. It will try to fetch the leader from the other brokers list which it has cached to check if there is any leader for this topic! As it does not find any, it will try to hit to the same leader (that is dead) and after reaching a maximum no of retries it will throw an exception !!

12. How replication works in kafka?

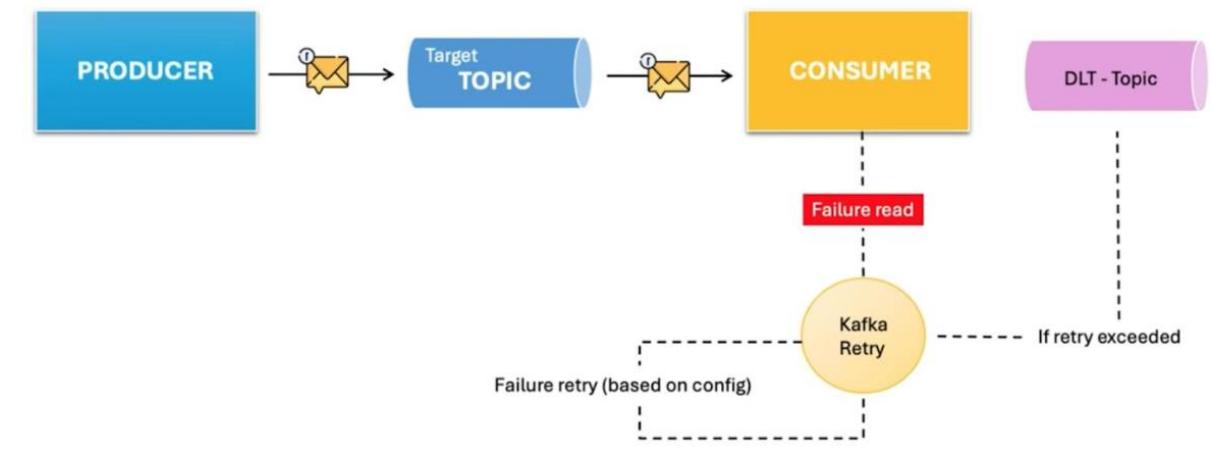
[Kafka Topic Replication - javatpoint](#)

[Replication in Kafka. Replication is the process of having... | by Aman Arora | Medium](#)

13. What is the default retry count of Kafka?

3

14. How to handle failures in Kafka?



Consider a scenario if you are processing the financial transaction failed due to system temporary issue. This can be handled with the help of retry mechanism. It will retry to process the failed transaction based on our configuration. If it is configured for 3 times it will retry for 3 attempts. It will create 2 DLT topics n-1. In case transaction processing failed in 2 attempts the that transaction will move to DLT (Dead Letter) Topic.

DLT Topic: DLT means Dead Letter Topic. If consumer is unable to process the messages even after retry attempts Kafka will store all the failure transactions events into DLT topic.

References:

[Exception Handling in Apache Kafka - GeeksforGeeks](#)

[Error Handling Patterns in Kafka](#)

Example: [Robust Kafka Consumer Error Handling on a Spring Boot 3 Application | by Noah Hsu | Javarevisited | Medium](#)

15. Why Apache Kafka is so Fast?

[Why Apache Kafka is so Fast? - GeeksforGeeks](#)

16. Give brief about Topics, Partitions, and Offsets in Apache Kafka?

1. Topics

- A **Topic** is a logical channel where messages (data) are published.
 - Producers send data to a specific topic, and consumers read from it.
 - Example: A topic named payment-events could store all payment-related data.
-

2. Partitions

- **Partitions** are sub-divisions of a topic that allow parallelism and scalability.
- Each topic is split into multiple **partitions**, and each partition stores a subset of the topic's data.
- Data within a partition is **ordered**, but there is no global order across partitions.
- Example: If a topic user-logs has 3 partitions, the data will be distributed across these partitions (P0, P1, P2).

Benefits:

- Parallelism: Multiple consumers can read from different partitions simultaneously.
 - Scalability: More partitions mean higher throughput.
-

3. Offsets

- An **Offset** is a unique identifier for each message within a partition.
- It is a sequential number assigned to each message as it gets written to a partition.
- Offsets are **immutable** and persist until data retention policies delete them.

Key Concept:

- Consumers use offsets to keep track of their read position.
 - Example: If a consumer has processed up to offset 50, it will resume from 51 next time.
-

Summary Relationship:

- **Topic** = Collection of **Partitions**.
- **Partition** = Ordered sequence of messages identified by **Offsets**.

[Topics, Partitions, and Offsets in Apache Kafka - GeeksforGeeks](#)

17. Explain possible Kafka failures?

Kafka failures can occur at various levels of the system, affecting the performance and reliability of your distributed messaging setup. Here are some common Kafka failure types and causes:

1. Broker Failures

Crash or Stop: A broker might stop due to resource exhaustion, hardware failure, or misconfiguration.

Network Partitioning: Brokers might lose connectivity, leading to data inconsistency or unavailability.

Disk Failures: Kafka relies heavily on disk storage. If a disk fails, it can result in data loss or unavailability.

2. Producer Failures

Message Loss: If `acks=0` or `acks=1` is used, messages can be lost before they are fully committed.

Timeouts: Network latency or broker unavailability can cause producer timeouts.

Message Duplication: In case of retries without idempotency enabled, messages might be duplicated.

3. Consumer Failures

Consumer Lag: If consumers fall behind in reading messages, it can result in high lag and delayed processing.

Consumer Group Rebalancing: Frequent rebalances can disrupt message processing.

Offset Management Issues: Incorrect handling of offsets can lead to reprocessing or message loss.

4. Data Loss and Inconsistency

Replication Issues: If replication is misconfigured or insufficient (e.g., replication factor < 3), data loss can occur.

Unclean Leader Elections: If `unclean.leader.election.enable` is true, data loss can occur when a non-synchronized replica becomes the leader.

5. Zookeeper Failures (for older Kafka versions)

Zookeeper Quorum Loss: If a majority of Zookeeper nodes fail, Kafka might become unavailable.

Leader Election Failure: Zookeeper issues can prevent leader elections from completing.

6. Configuration Errors

Retention Misconfiguration: Data might be deleted too soon if retention periods are too short.

Resource Mismanagement: Insufficient disk space, memory, or CPU can degrade Kafka's performance.

7. Mitigation Strategies

Replication and Acknowledgment: Set `acks=all` and `min.insync.replicas` appropriately.

Monitoring and Alerting: Use tools like Prometheus and Grafana for Kafka metrics monitoring.

Fault Tolerance: Ensure proper replication factors and redundancy.

Proper Partitioning: Distribute partitions evenly across brokers. Would you like help diagnosing a specific Kafka issue or configuring failure recovery?

18. Explain about Kafka internal data workflow?

Kafka's internal workflow involves the movement of data from producers to brokers and then to consumers, ensuring fault tolerance, scalability, and high throughput. Here's a step-by-step breakdown of Kafka's internal workflow:

1. Producer Workflow (Data Ingestion)

- Producers are client applications that send messages to Kafka.
- Steps:
 - Producer connects to a Kafka broker.
 - Producer sends messages to a topic.
 - If a key is provided, Kafka hashes the key to determine the partition.
 - If no key is provided, Kafka distributes messages in a round-robin fashion across partitions.

 Key Concept: Producers can send messages with acks=all for guaranteed durability by waiting for all replicas to acknowledge the message.

2. Partitioning and Storage (Broker Side)

- Each topic is divided into partitions.
- Each partition is stored across one or more brokers.
- Kafka writes data in an append-only log for each partition.

Leader-Follower Model:

- Each partition has a leader broker responsible for write and read operations.
- Other brokers maintain follower replicas.

 Key Concept: Messages are written sequentially to disk, making Kafka extremely fast due to optimized disk usage and the OS page cache.

3. Replication and Fault Tolerance

- Kafka ensures fault tolerance through replication.
- Each partition has multiple replicas based on the replication factor.
- Leader Broker: Handles all read and write operations.
- Follower Brokers: Sync data from the leader.

 ISR (In-Sync Replicas):

- Brokers that are fully synchronized with the leader.

- If the leader fails, one of the ISR replicas is elected as the new leader.
-

4. Consumer Workflow (Data Consumption)

- Consumers read messages from partitions.
- Consumer Groups:
 - Consumers subscribe to a topic as part of a consumer group.
 - Each partition is consumed by only one consumer in the group to avoid duplicate processing.

 Key Concept: Consumers pull data from brokers, controlling the rate of consumption.

5. Offset Management (Tracking Read Position)

- Offsets represent the position of a message in a partition.
- Kafka stores offsets per consumer group to keep track of consumption progress.
- Offset Storage:
 - Kafka stores offsets in an internal topic called __consumer_offsets.
 - Consumers can use auto-commit or manage offsets manually.

 Key Concept: Offsets allow consumers to resume from the last read message after a restart.

6. Zookeeper / KRaft (Cluster Management)

- Zookeeper (Older Versions)
 - Manages broker metadata, leader election, and health monitoring.
- KRaft Mode (Kafka 2.8+)
 - Kafka now offers a Zookeeper-free architecture with KRaft for metadata management and leader election.

7. Data Retention and Cleanup

- Kafka retains data for a configured time or size, even after being consumed.
- Retention Policies:
 - Time-based: Data is deleted after a set retention period.
 - Size-based: Data is deleted when the partition size exceeds a threshold.
 - Log Compaction: Keeps only the latest message per key for a partition.

Summary of Kafka Internal Workflow:

1. Producer sends messages to a Kafka topic.
2. Broker (Leader Partition) receives the message and writes it to disk.
3. Follower Replicas replicate the message for fault tolerance.
4. Consumer reads messages from partitions and commits offsets.
5. Zookeeper/KRaft manages metadata, leader elections, and cluster health.

19. How to handle and recover brokers failure?

Kafka is designed for fault tolerance and high availability. When a broker fails, Kafka can continue functioning without data loss or downtime due to its replication mechanism and leader election process.

1. Broker Failure Scenario:

A Kafka broker failure occurs when a broker goes offline, either due to hardware issues, network failures, or system crashes. Consequences include:

The broker stops handling read/write operations for its partitions.

Partitions handled by the broker may become unavailable temporarily.

2. How Kafka Handles Broker Failures Automatically:

Step 1: Leader Detection and Failure Identification

Each partition has a leader broker responsible for reads and writes.

Kafka uses Zookeeper (or KRaft in newer versions) to monitor broker health.

If a broker fails, Zookeeper detects the failure and notifies the cluster.

Step 2: Leader Election (Failover Mechanism)

If the leader broker fails, Kafka elects a new leader from the ISR (In-Sync Replicas).

The next available in-sync replica (ISR) becomes the new leader.

If no ISR is available, Kafka may perform an unclean leader election (if enabled).

Step 3: Client Redirection

Once a new leader is elected:

Producers and consumers automatically detect the new leader.

Kafka clients reconnect to the new leader broker without manual intervention.

Step 4: Replication Catch-Up (When Broker Recovers)

When the failed broker recovers:

It attempts to rejoin the cluster.

Kafka will synchronize its data with the current leader replica.

Once fully synced, it rejoins the ISR.

3. Key Configurations for Better Fault Tolerance:

a. Replication Factor:

Set replication.factor=3 (at least 3 replicas for redundancy).

b. Minimum In-Sync Replicas (min.insync.replicas):

Ensures a minimum number of replicas must acknowledge writes before success.

Recommended: min.insync.replicas=2.

c. Acknowledgment Settings (acks):

Producer setting: acks=all ensures writes are acknowledged only after all replicas sync.

d. Unclean Leader Election (unclean.leader.election.enable):

Disabled by default (false) to avoid data loss.

Enabling it (true) can risk data loss but allows faster recovery if no ISR is available.

4. Manual Broker Recovery Steps (If Automatic Fails):

Identify the Failed Broker:

Use **kafka-topics.sh** and **zookeeper-shell.sh** to inspect broker status.

Restart the Broker:

Ensure the broker configuration (server.properties) is correct.

Verify Replicas and ISR:

Check if the broker has rejoined the ISR using:

```
kafka-topics.sh --describe --topic <topic-name> --bootstrap-server <broker-list>
```

Reassign Partitions (If Needed):

If partitions remain under-replicated, reassign them using:

```
kafka-reassign-partitions.sh
```

5. Best Practices for Handling Broker Failures:

 **Replication:** Use a replication factor ≥ 3 .

 **Monitoring:** Use tools like Prometheus, Grafana, or Confluent Control Center for monitoring broker health.

- Avoid Unclean Leader Elections: Disable unclean leader election (false).
- Automated Recovery: Use Kubernetes or cluster management tools for auto-scaling and healing.