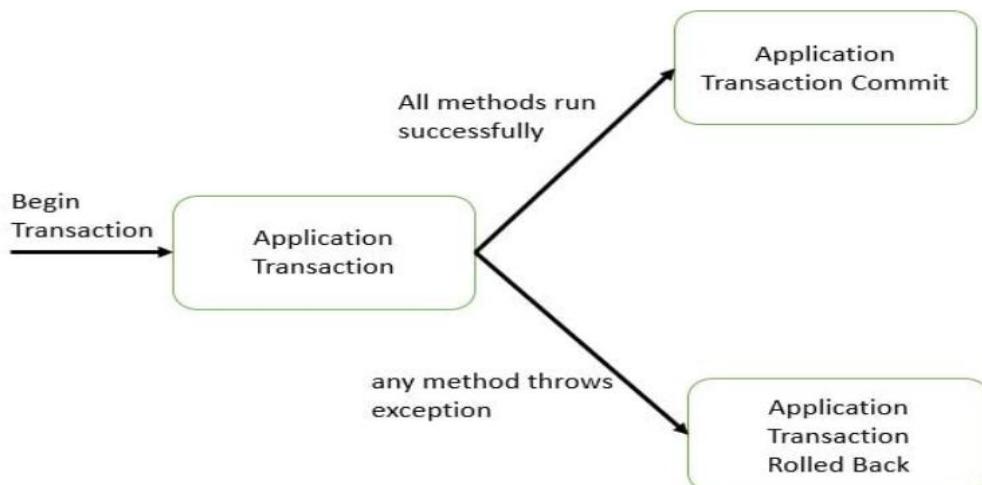


## 1. Explain Spring Transaction Management?

Spring provides powerful **transaction management** features that ensure data consistency and integrity in **multi-step operations** within applications. It handles transactions declaratively or programmatically, ensuring **ACID (Atomicity, Consistency, Isolation, Durability)** compliance in database operations.

Let's understand transactions with the above example. If a user has entered his information, the user's information gets stored in the user\_info table. Now, to book a ticket, he makes an online payment, and due to some reason(system failure), the payment has been canceled, so the ticket is not booked for him. But, the problem is that his information gets stored on the user\_info table. On a large scale, more than thousands of these things happen within a single day. So, it is not good practice to store a single action of the transaction(Here, only user info is stored, not the payment info).

To overcome these problems, Spring provides transaction management, which uses annotation to handle these issues. In such a scenario, the spring stores the user information in temporary memory and then checks for payment information. If the payment is successful, then it will complete the transaction; otherwise, it will roll back the transaction, and the user information will not be stored in the database.



## 2. How to handle the Spring Transaction Management?

### @Transactional Annotation

In Spring Boot, `@Transactional` annotation is used to manage transactions in a Spring Boot application and to define a scope of transaction. This annotation can be applied to the class level or method level. It provides data reliability and consistency. When a method is indicated with `@Transactional` annotation, it indicates that the particular method should be executed within the context of that transaction. If the transaction becomes successful, then the changes made to the database are committed; if any transaction fails, all the changes made to that particular transaction can be rollback and it will ensure that the database remains in a consistent state.

*Note: To use `@Transactional` annotation, you need to configure transaction management. However, if you are using `spring-boot-starter-data-jpa`, Spring Boot auto-configures transaction management, and you do not need to explicitly add `@EnableTransactionManagement`. If you are not using `spring-boot-starter-data-jpa` or need custom transaction management, you can enable it by adding `@EnableTransactionManagement` to your main class of the Spring Boot application.*

### 3. Explain different types of Transaction management?

#### Types of Spring Transactions

##### 1. Declarative Transaction Management (Recommended)

- Uses `@Transactional` annotation on **service-layer methods**.
- No need for explicit `beginTransaction()` or `commit()` calls.

##### 2. Programmatic Transaction Management

- Uses `TransactionTemplate` or `PlatformTransactionManager` to **manually control transactions**.
- Offers flexibility for **custom rollback** and commit handling.

#### Example of Declarative Transaction Management

Java

 Copy

```
import org.springframework.transaction.annotation.Transactional;
import org.springframework.stereotype.Service;

@Service
public class BankService {

    @Transactional
    public void transferMoney(Account from, Account to, double amount) {
        from.withdraw(amount);
        to.deposit(amount);
    }
}
```

If an error occurs during withdrawal or deposit, **Spring will automatically roll back the transaction**.

#### Example of Programmatic Transaction Management

```
import org.springframework.transaction.PlatformTransactionManager;
import org.springframework.transaction.TransactionDefinition;
import org.springframework.transaction.TransactionStatus;
import org.springframework.transaction.support.DefaultTransactionDefinition;

public class TransactionalService {
    private PlatformTransactionManager transactionManager;

    public TransactionalService(PlatformTransactionManager transactionManager) {
        this.transactionManager = transactionManager;
    }

    public void processTransaction() {
        TransactionDefinition def = new DefaultTransactionDefinition();
        TransactionStatus status = transactionManager.getTransaction(def);

        try {
            // Perform database operations
            transactionManager.commit(status);
        } catch (Exception e) {
            transactionManager.rollback(status);
        }
    }
}
```

#### 4. Give an example for Spring Boot Declarative transaction management?

##### Step 1: Create A Spring Boot Project

In this step, we will create a spring boot project. For this, we will be using [Spring Initializr](#). To create a spring boot project please refer to [How to Create a Spring Boot Project?](#)

##### Step 2: Add Dependencies

We will add the required dependencies for our spring boot application.

The screenshot shows the Spring Initializr interface. In the 'Project' section, 'Maven' is selected. In the 'Language' section, 'Java' is selected. In the 'Spring Boot' section, version '3.0.1' is selected. Under 'Project Metadata', the group is set to 'com.geeksforgeeks', artifact to 'transaction-management', name to 'transaction-management', description to 'Demo project for Spring Boot', package name to 'com.geeksforgeeks.transaction-management', and packaging to 'Jar'. In the 'Dependencies' section, several tools are listed: 'Spring Boot DevTools' (Developer Tools), 'Lombok' (Developer Tools), 'Spring Web' (Web), 'Spring Data JPA' (SQL), and 'MySQL Driver' (SQL).

##### Step 3: Configure Database

Now, we will configure the database in our application. We will be using the following configurations and add them to our `application.properties` file.

```
server.port = 9090
#database configuration
spring.datasource.url=jdbc:mysql://localhost:3306/employee_db
spring.datasource.username=root
spring.datasource.password=root
spring.jpa.properties.hibernate.dialect=org.hibernate.dialect.MySQL57Dialect
#the ddl-auto:update : It will create the entity schema and map it to db automatically
spring.jpa.hibernate.ddl-auto=update
spring.jpa.show-sql=true
```

*Note: Please add your database username & password along with the database path.*

#### Step 4: Create Model Class

In this step, we will create our model class. Here, we will be creating two model classes, **Employee** and **Address**. While creating the model class we will be using [Lombok Library](#).

**Employee.java**

```
import jakarta.persistence.Entity;
import jakarta.persistence.GeneratedValue;
import jakarta.persistence.GenerationType;
import jakarta.persistence.Id;
import jakarta.persistence.Table;
import lombok.AllArgsConstructor;
import lombok.Getter;
import lombok.NoArgsConstructor;
import lombok.Setter;
import lombok.ToString;

@Getter
@Setter
@NoArgsConstructor
@AllArgsConstructor
@ToString
@Entity
@Table(name="EMP_INFO")
public class Employee {

    @Id
    @GeneratedValue(strategy = GenerationType.AUTO)
    private int id;
    private String name;
}
```

**Address.java**

```
import jakarta.persistence.Entity;
import jakarta.persistence.GeneratedValue;
import jakarta.persistence.GenerationType;
import jakarta.persistence.Id;
import jakarta.persistence.OneToOne;
import jakarta.persistence.Table;
import lombok.AllArgsConstructor;
import lombok.Getter;
import lombok.NoArgsConstructor;
import lombok.Setter;
import lombok.ToString;

@Getter
@Setter
@NoArgsConstructor
@AllArgsConstructor
@ToString
@Entity
@Table(name="ADD_INFO")
public class Address {

    @Id
    @GeneratedValue(strategy = GenerationType.AUTO)
    private Long id;
    private String address;

    // one to one mapping means,
    // one employee stays at one address only
    @OneToOne
    private Employee employee;

    @Version
    private Long version;
}
```

---

## Step 5: Create a Database Layer

In this step, we will create a database layer. For this, we will be creating **EmployeeRepository** and **AddressRepository** and will be extending **JpaRepository<T, ID>** for performing database-related queries.

### EmployeeRepository.java

```
import org.springframework.data.jpa.repository.JpaRepository;
import com.geeksforgeeks.transactionmanagement.model.Employee;

public interface EmployeeRepository extends JpaRepository<Employee, Integer> {
}
```

### AddressRepository.java

```
import org.springframework.data.jpa.repository.JpaRepository;
import com.geeksforgeeks.transactionmanagement.model.Address;

public interface AddressRepository extends JpaRepository<Address, Integer> {
}
```

## Step 6: Create a Service Layer

You can use **@Transactional** annotation in service layer which will result interacting with the database. In this step, we will create a service layer for our application and add business logic to it. For this, we will be creating two classes **EmployeeService** and **AddressService**. In **EmployeeService** class we are throwing an exception.

### EmployeeService.java

```
import org.springframework.beans.factory.annotation.Autowired;
import org.springframework.stereotype.Service;
import org.springframework.transaction.annotation.Transactional;

import com.geeksforgeeks.transactionmanagement.model.Address;
import com.geeksforgeeks.transactionmanagement.model.Employee;
import com.geeksforgeeks.transactionmanagement.repository.EmployeeRepository;

@Service
public class EmployeeService {

    @Autowired
    private EmployeeRepository employeeRepository;

    @Autowired
    private AddressService addressService;

    @Transactional
    public Employee addEmployee(Employee employee) throws Exception {
        Employee employeeSavedToDB = this.employeeRepository.save(employee);

        Address address = new Address();
        address.setId(123L);
        address.setAddress("Varanasi");
        address.setEmployee(employee);

        // calling addAddress() method
        // of AddressService class
        this.addressService.addAddress(address);
        return employeeSavedToDB;
    }
}
```

## AddressService.java

```
import org.springframework.beans.factory.annotation.Autowired;
import org.springframework.stereotype.Service;

import com.geeksforgeeks.transactionmanagement.model.Address;
import com.geeksforgeeks.transactionmanagement.repository.AddressRepository;

@Service
public class AddressService {

    @Autowired
    private AddressRepository addressRepository;

    public Address addAddress(Address address) {
        Address addressSavedToDB = this.addressRepository.save(address);
        return addressSavedToDB;
    }

}
```

### Step 7: Create Controller

In this step, we will create a controller for our application. For this, we will create a Controller class and add all the mappings to it.

#### Controller.java

```
import org.springframework.beans.factory.annotation.Autowired;
import org.springframework.http.HttpStatus;
import org.springframework.http.ResponseEntity;
import org.springframework.web.bind.annotation.PostMapping;
import org.springframework.web.bind.annotation.RequestBody;
import org.springframework.web.bind.annotation.RequestMapping;
import org.springframework.web.bind.annotation.RestController;

import com.geeksforgeeks.transactionmanagement.model.Employee;
import com.geeksforgeeks.transactionmanagement.service.EmployeeService;

@RestController
@RequestMapping("/api/employee")
public class Controller {

    @Autowired
    private EmployeeService employeeService;

    @PostMapping("/add")
    public ResponseEntity<Employee> saveEmployee(@RequestBody Employee employee) throws Exception{
        Employee employeeSavedToDB = this.employeeService.addEmployee(employee);
        return new ResponseEntity<Employee>(employeeSavedToDB, HttpStatus.CREATED);
    }
}
```

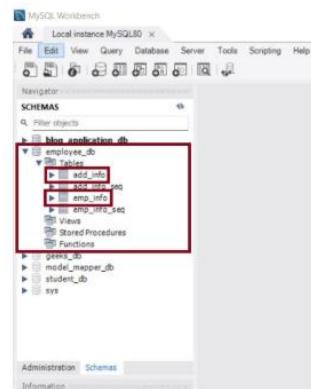
### Step 8: Running Our Application

In this step, we will run our application. Once, we run our application using hibernate mapping in our database required tables will be created.



```
2023-01-15T14:10:00.359+05:30  WARN 12048 --- [ restartedMain] org.hibernate.orm.deprecation : HHH00000021: Encountered deprecate
2023-01-15T14:10:00.531+05:30  INFO 12048 --- [ restartedMain] com.zaxxer.hikari.HikariDataSource : HikariPool-1 - Starting...
2023-01-15T14:10:01.013+05:30  INFO 12048 --- [ restartedMain] com.zaxxer.hikari.pool.HikariPool : HikariPool-1 - Added connection c
2023-01-15T14:10:01.016+05:30  INFO 12048 --- [ restartedMain] com.zaxxer.hikari.HikariDataSource : HikariPool-1 - Start completed.
2023-01-15T14:10:01.048+05:30  INFO 12048 --- [ restartedMain] org.hibernate.orm.deprecation : HHH000400: Using dialect: org.hib
2023-01-15T14:10:01.049+05:30  WARN 12048 --- [ restartedMain] org.hibernate.orm.deprecation : HHH0000025: MySQLSelect has b
Hibernate: create table add_info (id bigint not null, address varchar(255), employee_id integer, primary key (id)) engine=InnoDB
Hibernate: create table add_info_seq (next_val bigint) engine=InnoDB
Hibernate: insert into add_info_seq values ( 1 )
Hibernate: create table emp_info (id integer not null, name varchar(255), primary key (id)) engine=InnoDB
Hibernate: create table emp_info_seq (next_val bigint) engine=InnoDB
Hibernate: insert into emp_info_seq values ( 1 )
Hibernate: alter table add_info add constraint FKM3gr2uljsypowcg8ljevhesi foreign key (employee_id) references emp_info (id)
2023-01-15T14:10:02.634+05:30  INFO 12048 --- [ restartedMain] o.h.e.t.j.p.l.JpaLocalContainerEntityManagerFactoryBean : Initialized JPA EntityManagerFact
2023-01-15T14:10:02.845+05:30  INFO 12048 --- [ restartedMain] j.LocalContainerEntityManagerFactoryBean : Initialized JPA EntityManagerFact
2023-01-15T14:10:03.303+05:30  WARN 12048 --- [ restartedMain] JpaBaseConfiguration$JpaWebConfiguration : spring.jpa.open-in-view is enable
2023-01-15T14:10:03.905+05:30  INFO 12048 --- [ restartedMain] o.s.b.d.a.OptionalLiveReloadServer : LiveReload server is running a
2023-01-15T14:10:03.993+05:30  INFO 12048 --- [ restartedMain] o.s.w.e.tomcat.TomcatWebServer : Tomcat started on port(s): 9090 ( 
2023-01-15T14:10:04.008+05:30  INFO 12048 --- [ restartedMain] c.g.t.TransactionManagementApplication : Started TransactionManagementAppl
```

As we can see in logs, our table has been created. We can also confirm it by looking at our database.



Now, we will request our application for adding an employee to it, using postman. To learn more about postman please refer to [Postman – Working, HTTP Request & Responses](#). Once, we hit the request, **the request moves from the controller to the service layer where our business logic is present.**

A screenshot of the Postman application. The request URL is 'http://localhost:9090/api/employees/add'. The method is 'POST'. The 'Body' tab shows a JSON payload: { "name": "ankur" }. The response status is 201 Created, and the JSON data returned is { "id": 1, "name": "ankur" }.

As we can see in the above response we have added an employee. We can also check our database for employee data and address data.

A screenshot of MySQL Workbench showing the 'emp\_info' table. A red box highlights the SQL query: 'SELECT \* FROM employee\_db.emp\_info;'. The result grid shows one row with id 1 and name 'ankur'.

Similarly, we can also check for address data.

A screenshot of MySQL Workbench showing the 'add\_info' table. A red box highlights the SQL query: 'SELECT \* FROM employee\_db.add\_info;'. The result grid shows one row with id 1 and address 'Varanasi'.

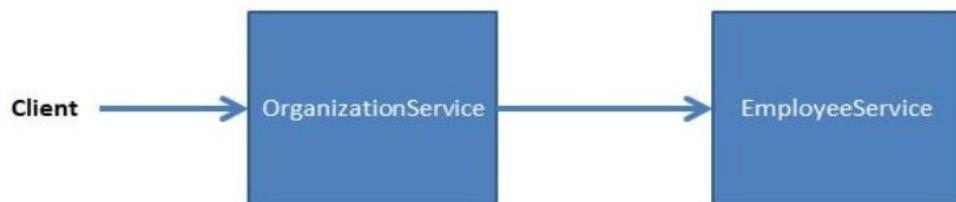
## 5. Explain different types of propagation in transaction management?

Transaction propagation determines how transactions interact when multiple transactional methods are executed. Spring provides different **propagation types** to define how transactions behave within nested calls.

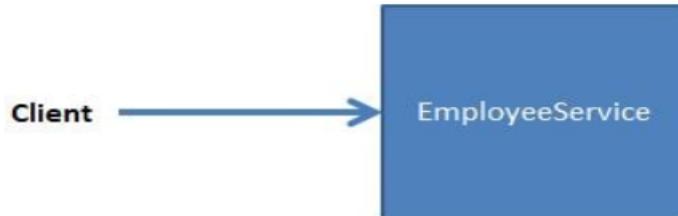
Propagation	Behaviour
REQUIRED	<b>Always executes in a transaction.</b> If there is any existing transaction it uses it. If none exists then only a new one is created
SUPPORTS	<b>It may or may not run in a transaction.</b> If current transaction exists then it is supported. If none exists then gets executed with out transaction.
NOT_SUPPORTED	<b>Always executes without a transaction.</b> If there is any existing transaction it gets suspended
REQUIRES_NEW	<b>Always executes in a new transaction.</b> If there is any existing transaction it gets suspended
NEVER	<b>Always executes with out any transaction.</b> It throws an exception if there is an existing transaction
MANDATORY	<b>Always executes in a transaction.</b> If there is any existing transaction it is used. If there is no existing transaction it will throw an exception.

But suppose the user wants to call the Employee Service in both ways i.e.

- Call using Organization service



- Call the the Employee Service directly.

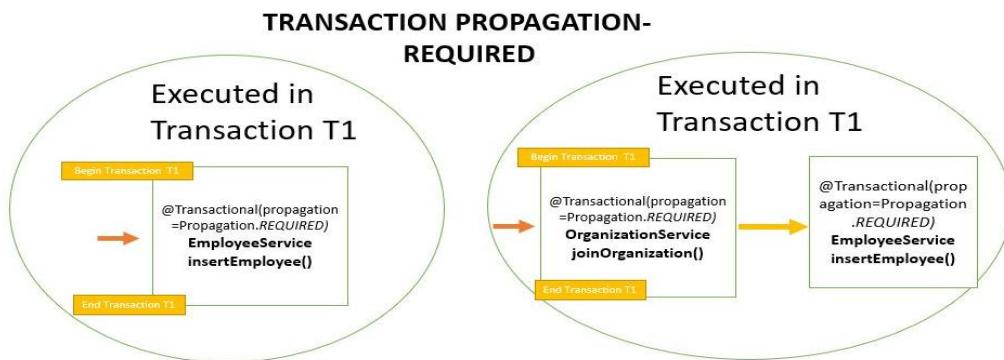


As the Employee Service may also be called directly we will need to use Transaction annotation with Employee Service also. So both the services - Organization Service and the Employee Service will be using Transaction annotation.

We will be looking at the various propagation scenarios by observing the behaviour of the Organization and Employee service. There are six types of Transaction Propagations-

- **REQUIRED**
- **SUPPORTS**
- **NOT\_SUPPORTED**
- **REQUIRES\_NEW**
- **NEVER**
- **MANDATORY**

### **Transaction Propagation - REQUIRED (Default Transaction Propagation)**



If the `insertEmployee()` method is called directly it creates its own new transaction.

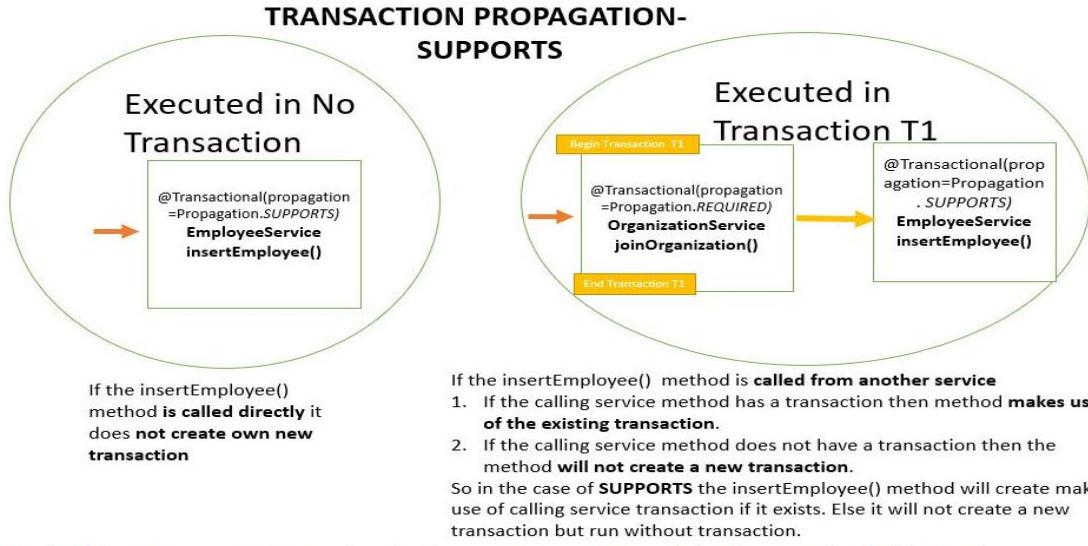
If the `insertEmployee()` method is called from another service –  
 1. If the calling service has a transaction then method makes use of the existing transaction.  
 2. If the calling service does not have a transaction then the method will create new transaction.  
 So in the case of REQUIRED the `insertEmployee()` method makes use of the calling service transaction if it exists else creates its own.

Here both the Organization Service and the Employee Service have the transaction propagation defined as **Required**. This is the default Transaction Propagation.

## Example

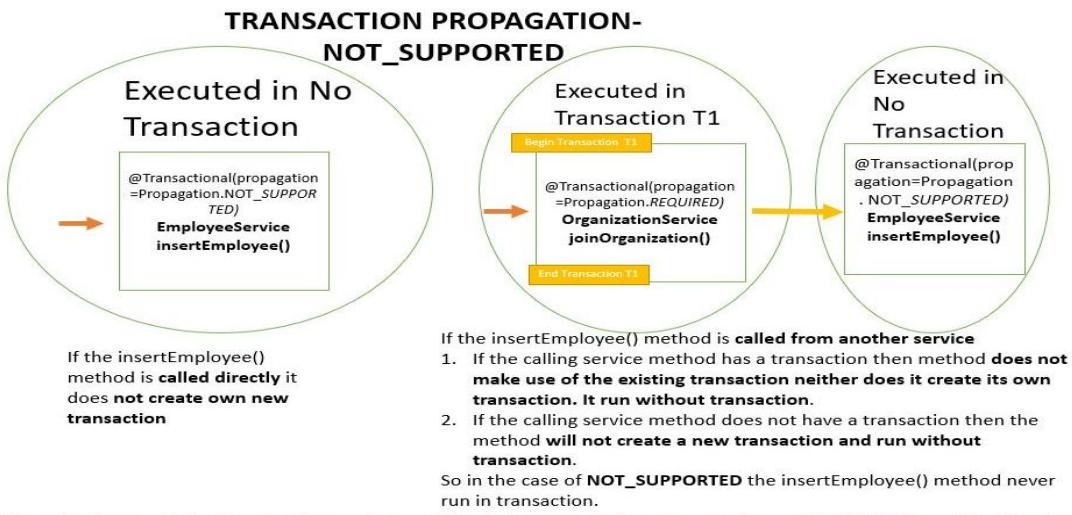
```
@Transactional(propagation = Propagation.REQUIRED)
public void requiredExample(String user) {
    // ...
}
```

## Transaction Propagation - SUPPORTS



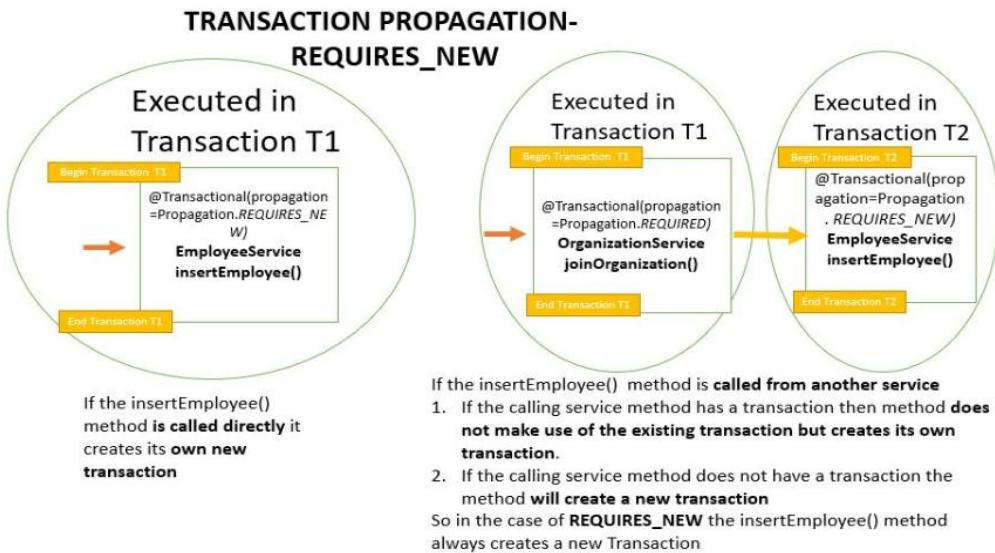
Here both the Organization Service has the transaction propagation defined as **Required** while Employee Service the transaction propagation is defined as **Supports**.

## Transaction Propagation - NOT\_SUPPORTED



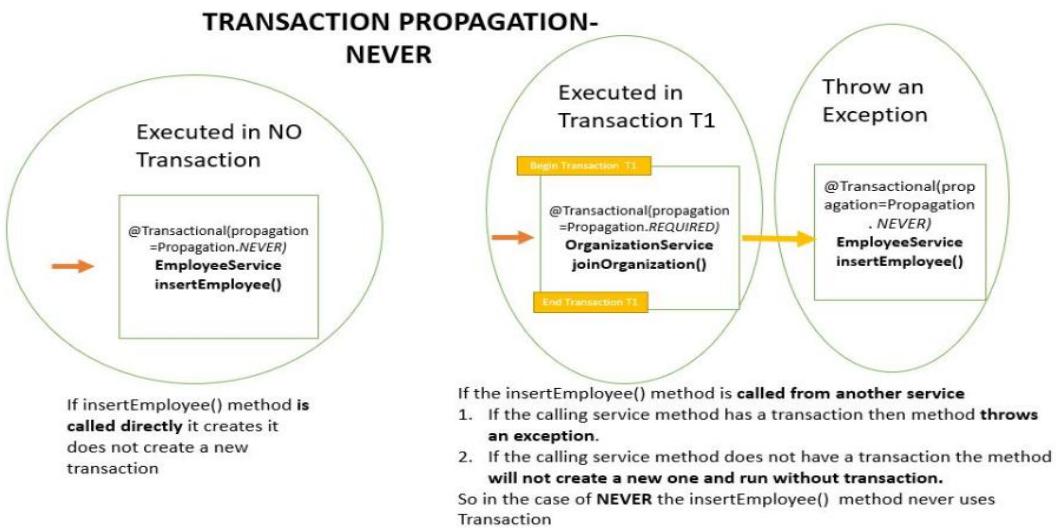
Here for the Organization Service we have defined the transaction propagation as **REQUIRED** and the Employee Service have the transaction propagation defined as **NOT\_SUPPORTED**

## Transaction Propagation - REQUIRES\_NEW



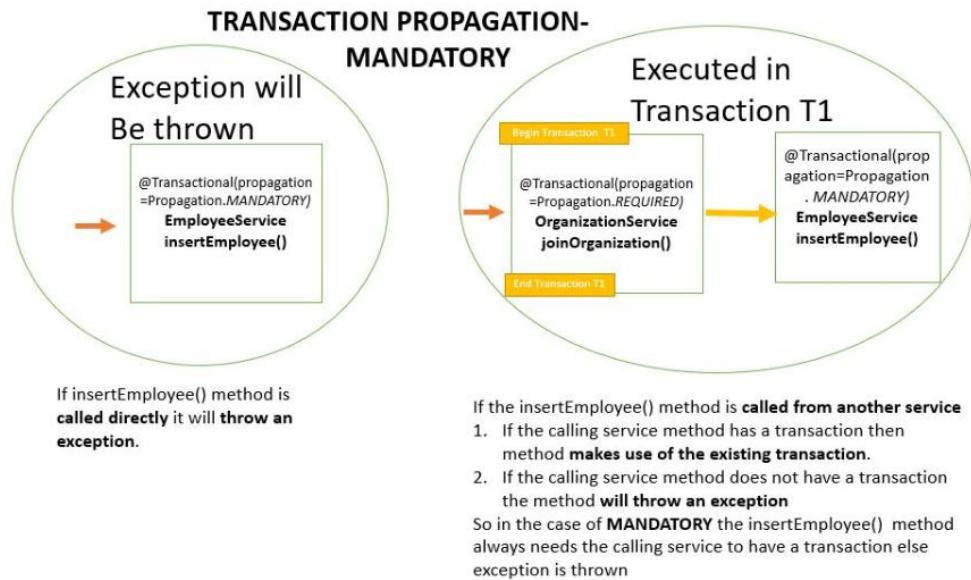
Here for the Organization Service we have defined the transaction propagation as **REQUIRED** and the Employee Service have the transaction propagation defined as **REQUIRES\_NEW**

## Transaction Propagation - NEVER



Here for the Organization Service we have defined the transaction propagation as **REQUIRED** and the Employee Service have the transaction propagation defined as **NEVER**

## Transaction Propagation - MANDATORY



Here for the Organization Service we have defined the transaction propagation as **REQUIRED** and the Employee Service have the transaction propagation defined as **MANDATORY**

## 6. Explain about different types of Transaction Isolation?

**Transaction Isolation** controls **how and when** the changes made by one transaction become visible to other concurrent transactions. It's a key part of maintaining **data consistency**.

### Types of Transaction Isolation Levels

Spring supports the following **isolation levels**, based on database configurations:

Isolation Level	Description
<b>READ UNCOMMITTED</b>	Transactions see <b>uncommitted changes</b> of other transactions, leading to <b>dirty reads</b> .
<b>READ COMMITTED</b>	Transactions see <b>only committed changes</b> , preventing <b>dirty reads</b> but allowing <b>non-repeatable reads</b> .
<b>REPEATABLE READ</b>	Ensures consistent reads, preventing <b>dirty and non-repeatable reads</b> , but not <b>phantom reads</b> .
<b>SERIALIZABLE</b>	The strictest isolation, completely <b>avoiding concurrency issues</b> but reducing performance.

## SERIALIZABLE

If two transactions are executing concurrently then it is as if the transactions get executed serially i.e the first transaction gets committed only then the second transaction gets executed. This is **total isolation**. So a running transaction is never affected by other transactions. However this may cause issues as **performance will be low and deadlock might occur**.

Transaction1	Transaction2
<p>Begin Transaction1. On select we get 3 records</p> <pre>mysql&gt; SET SESSION TRANSACTION ISOLATION LEVEL SERIALIZABLE; Query OK, 0 rows affected (0.00 sec)  mysql&gt; SELECT @@TX_ISOLATION; +-----+   @@TX_ISOLATION   +-----+   SERIALIZABLE      +-----+ 1 row in set, 1 warning (0.00 sec)  mysql&gt; begin; Query OK, 0 rows affected (0.00 sec)  mysql&gt; select * from employee; +-----+-----+   empId   empName   +-----+-----+   emp1    emp1       emp2    emp2       emp3    emp3     +-----+-----+ 3 rows in set (0.00 sec)</pre>	
	<p>Begin Transaction2. Try inserting a record. It will not allow to insert</p> <pre>mysql&gt; SET SESSION TRANSACTION ISOLATION LEVEL SERIALIZABLE; Query OK, 0 rows affected (0.00 sec)  mysql&gt; SELECT @@TX_ISOLATION; +-----+   @@TX_ISOLATION   +-----+   SERIALIZABLE      +-----+ 1 row in set, 1 warning (0.00 sec)  mysql&gt; set autocommit=0; Query OK, 0 rows affected (0.00 sec)  mysql&gt; begin; Query OK, 0 rows affected (0.00 sec)  mysql&gt; insert into employee(empId,empName)values('emp4','emp4'); ERROR 1205 (HY000): Lock wait timeout exceeded; try restarting transaction mysql&gt;</pre>
<p>On doing a select again we fetch 3 records. The transaction2 is not allowed to insert a new record till transaction1 completes.</p> <pre>mysql&gt; select * from employee; +-----+-----+   empId   empName   +-----+-----+   emp1    emp1       emp2    emp2       emp3    emp3     +-----+-----+ 3 rows in set (0.00 sec)</pre>	

## REPEATABLE\_READ

If two transactions are executing concurrently - **till the first transaction is committed the existing records cannot be changed by second transaction but new records can be added**. After the second transaction is committed, the new added records get reflected in first transaction which is still not committed. For MySQL the default isolation level is REPEATABLE\_READ.

However the REPEATABLE READ isolation level behaves differently when using mysql. When using MYSQL we are not able to see the newly added records that are committed by the second transaction.

Transaction1	Transaction2
<p>Begin Transaction1. On select we get 2 records</p> <pre>mysql&gt; SELECT @@TX_ISOLATION; +-----+   @@@TX_ISOLATION   +-----+   REPEATABLE-READ   +-----+ 1 row in set, 1 warning (0.00 sec)  mysql&gt; set autocommit=0; Query OK, 0 rows affected (0.00 sec)  mysql&gt; BEGIN; Query OK, 0 rows affected (0.00 sec)  mysql&gt; select * from employee; +-----+   empId   empName   +-----+   emp1    emp1       emp2    emp2     +-----+ 2 rows in set (0.00 sec)</pre>	
	<p>Begin Transaction2. Insert a record and do commit.</p> <pre>mysql&gt; set autocommit=0; Query OK, 0 rows affected (0.00 sec)  mysql&gt; SELECT @@TX_ISOLATION; +-----+   @@@TX_ISOLATION   +-----+   REPEATABLE-READ   +-----+ 1 row in set, 1 warning (0.00 sec)  mysql&gt; begin -&gt;; Query OK, 0 rows affected (0.00 sec)  mysql&gt; insert into employee(empId,empName)values('emp3','emp3') Query OK, 1 row affected (0.00 sec)  mysql&gt; commit; Query OK, 0 rows affected (0.01 sec)</pre>
<p>On doing a select again we fetch 2 records. The record inserted in transaction2 is not reflected</p> <pre>mysql&gt; select * from employee; +-----+   empId   empName   +-----+   emp1    emp1       emp2    emp2     +-----+ 2 rows in set (0.00 sec)</pre>	

## READ\_COMMITTED

If two transactions are executing concurrently - **before the first transaction is committed the existing records can be changed as well as new records can be changed by second transaction.** After the second transaction is committed, the newly added and also updated records get reflected in first transaction which is still not committed.

Transaction1	Transaction2
<p>Begin Transaction1. On select we get 4 records</p> <pre>mysql&gt; SET SESSION TRANSACTION ISOLATION LEVEL READ COMMITTED; Query OK, 0 rows affected (0.00 sec)  mysql&gt; set autocommit=0; Query OK, 0 rows affected (0.00 sec)  mysql&gt; begin; Query OK, 0 rows affected (0.00 sec)  mysql&gt; select * from employee; +-----+-----+   empId   empName   +-----+-----+   emp1    emp1       emp2    emp2       emp3    emp3       emp4    emp4     +-----+-----+ 4 rows in set (0.00 sec)</pre>	
	<p>Begin Transaction2. Insert a record and do commit.</p> <pre>mysql&gt; SET SESSION TRANSACTION ISOLATION LEVEL READ COMMITTED; Query OK, 0 rows affected (0.00 sec)  mysql&gt; set autocommit=0; Query OK, 0 rows affected (0.00 sec)  mysql&gt; begin; Query OK, 0 rows affected (0.00 sec)  mysql&gt; insert into employee(empId,empName)values('emp5','emp5'); Query OK, 1 row affected (0.00 sec)  mysql&gt; commit; Query OK, 0 rows affected (0.01 sec)</pre>
<p>On doing a select again we fetch 5 records. The record inserted in transaction2 is reflected in transaction1 only after transaction 2 is committed.</p> <pre>mysql&gt; select * from employee; +-----+-----+   empId   empName   +-----+-----+   emp1    emp1       emp2    emp2       emp3    emp3       emp4    emp4       emp5    emp5     +-----+-----+ 5 rows in set (0.00 sec)</pre>	

## READ\_UNCOMMITTED

If two transactions are executing concurrently - before the first transaction is committed the existing records can be changed as well as new records can be changed by second transaction. **Even if the second transaction is not committed the newly added and also updated records get reflected** in first transaction which is still not committed.

Transaction1	Transaction2
<p><b>Begin Transaction1. On select we get 3 records</b></p> <pre>mysql&gt; SET SESSION TRANSACTION ISOLATION LEVEL READ UNCOMMITTED; Query OK, 0 rows affected (0.00 sec)  mysql&gt; set autocommit=0; Query OK, 0 rows affected (0.00 sec)  mysql&gt; begin; Query OK, 0 rows affected (0.00 sec)  mysql&gt; select * from employee; +-----+-----+   empId   empName   +-----+-----+   emp1    emp1       emp2    emp2       emp3    emp3     +-----+-----+ 3 rows in set (0.00 sec)</pre>	
<p><b>On doing a select again we fetch 4 records. The record inserted in transaction2 is not committed, but still gets reflected in transaction1.</b></p> <pre>mysql&gt; select * from employee; +-----+-----+   empId   empName   +-----+-----+   emp1    emp1       emp2    emp2       emp3    emp3       emp4    emp4     +-----+-----+ 4 rows in set (0.00 sec)</pre>	<p><b>Begin Transaction2. Insert a record and do not do a commit.</b></p> <pre>mysql&gt; SET SESSION TRANSACTION ISOLATION LEVEL READ UNCOMMITTED; Query OK, 0 rows affected (0.00 sec)  mysql&gt; set autocommit=0; Query OK, 0 rows affected (0.00 sec)  mysql&gt; SELECT @@TX_ISOLATION; +-----+   @@TX_ISOLATION   +-----+   READ-UNCOMMITTED   +-----+ 1 row in set, 1 warning (0.00 sec)  mysql&gt; begin; Query OK, 0 rows affected (0.00 sec)  mysql&gt; insert into employee(empId,empName)values('emp4','emp4'); Query OK, 1 row affected (0.00 sec)</pre>

### Summary

- **Dirty Reads** - Suppose two transactions - Transaction A and Transaction B are running concurrently. If Transaction A modifies a record but not commits it. Transaction B reads this record but then Transaction A again rollbacks the changes for the record and commits it. So Transaction B has a wrong value.
- **Non-Repeatable Reads** - Suppose two transactions - Transaction A and Transaction B are running concurrently. If Transaction A reads some records. Transaction B modifies these records before transaction A has been committed. So if Transaction A again reads these records they will be different. So same select statements result in different existing records.
- **Phantom Reads** - Suppose two transactions - Transaction A and Transaction B are running concurrently. If Transaction A reads some records. Transaction B adds more such records before transaction A has been committed. So if Transaction A again reads there will be more records than the previous select statement. So same select statements result in different number records to be displayed as new records also get added.

Isolation Level	Dirty Reads	Non-Repeatable Reads	Phantom Reads
SERIALIZABLE	This scenario is not possible as the second transaction cannot start execution until the first is committed. They never execute parallelly but only sequentially	This scenario is not possible as the second transaction cannot start execution until the first is committed. They never execute parallelly but only sequentially	This scenario is not possible as the second transaction cannot start execution until the first is committed. They never execute parallelly but only sequentially
REPEATABLE_READ	This scenario is not possible as any existing record change gets reflected only if the transaction is committed. So other transaction will never read wrong value.	This scenario is not possible since any record can be changed only after a transaction has been committed. So multiple select statements before transaction commit will always return same existing records.	This scenario is possible as other transactions can insert new records even if first transaction commit has not taken place.
READ_COMMITTED	This scenario is not possible as any existing record change gets reflected only if the transaction is committed. So other transaction will never read wrong value.	This scenario is possible as other transactions can modify existing records even if first transaction commit has not taken place.	This scenario is possible as other transactions can insert new records even if first transaction commit has not taken place.
READ_UNCOMMITTED	This scenario is possible as any record can be read by other transactions even if the first transaction is not committed. So if first transaction rollbacks the record changes then other transactions will have wrong values	This scenario is possible since any record can be changed even if a transaction is not committed.	This scenario is possible as any record can be inserted even if a transaction is not committed.

## Example:

```

@Service
public class OrganizationServiceImpl implements OrganizationService {

    @Autowired
    EmployeeService employeeService;

    @Autowired
    HealthInsuranceService healthInsuranceService;

    @Override
    // Using Transactional annotation we can define any isolation level supported by the underlying database.
    @Transactional(isolation = Isolation.SERIALIZABLE)
    public void joinOrganization(Employee employee, EmployeeHealthInsurance employeeHealthInsurance) {
        employeeService.insertEmployee(employee);
        healthInsuranceService.registerEmployeeHealthInsurance(employeeHealthInsurance);
    }

    @Override
    @Transactional
    public void leaveOrganization(Employee employee, EmployeeHealthInsurance employeeHealthInsurance) {
        employeeService.deleteEmployeeById(employee.getEmpId());
        healthInsuranceService.deleteEmployeeHealthInsuranceById(employeeHealthInsurance.getEmpId());
    }
}

```

## 7. Explain Rollback Mechanism in Spring Boot Transaction Management?

Rollback ensures **data integrity** by undoing changes **when an error occurs** during a transaction. Spring Boot **automatically handles rollback** in transactions based on **exception types and configurations**

## How Rollback Works in Spring Boot

- ✓ If a transaction **fails**, Spring **reverts all database changes** made within that transaction.
- ✓ By default, **unchecked exceptions** (`RuntimeException`, `Error`) **trigger rollback**.
- ✓ **Checked exceptions** (`Exception`) **do NOT trigger rollback automatically**, but can be explicitly configured.

## Example: Automatic Rollback on RuntimeException

```
import org.springframework.transaction.annotation.Transactional;
import org.springframework.stereotype.Service;

@Service
public class BankService {

    @Transactional
    public void transferMoney(Account from, Account to, double amount) {
        from.withdraw(amount);
        to.deposit(amount);

        if (amount > 10000) { // Simulating an error
            throw new RuntimeException("Transaction amount exceeds limit!");
        }
    }
}
```

### Rollback for Checked Exceptions

To roll back **checked exceptions**, explicitly set `rollbackFor` in `@Transactional`:

```
Java Copy
@Transactional(rollbackFor = Exception.class)
public void processOrder(Order order) throws Exception {
    order.reserveStock();
    order.makePayment();

    if (!order.isValid()) {
        throw new Exception("Invalid order!");
    }
}
```

💡 Without `rollbackFor`, Spring will NOT roll back for checked exceptions.

## Handling Rollback Programmatically

```

import org.springframework.transaction.PlatformTransactionManager;
import org.springframework.transaction.TransactionStatus;
import org.springframework.transaction.support.DefaultTransactionDefinition;

public class PaymentService {
    private PlatformTransactionManager transactionManager;

    public PaymentService(PlatformTransactionManager transactionManager) {
        this.transactionManager = transactionManager;
    }

    public void processTransaction() {
        TransactionStatus status = transactionManager.getTransaction(new DefaultTransactionDefinition());

        try {
            // Perform database operations
            transactionManager.commit(status);
        } catch (Exception e) {
            transactionManager.rollback(status); // Manual rollback
        }
    }
}

```

## 8. Can we apply @Transactional in class level and method level?

Yes, **@Transactional** can be applied at both the **class level** and **method level** in Spring.

### Class-Level **@Transactional**

- When applied at the **class level**, it **affects all methods** inside the class **unless a method has its own transactional configuration**.
- Provides a **default transaction behavior** for the entire class.
- Individual methods **can override** the class-level transaction settings.

### Example: Applying **@Transactional** at Class Level

Java

 Copy

```

import org.springframework.transaction.annotation.Transactional;
import org.springframework.stereotype.Service;

@Service
@Transactional // Applied to all methods in this class
public class AccountService {

    public void deposit(Account account, double amount) {
        account.addBalance(amount);
    }

    public void withdraw(Account account, double amount) {
        account.reduceBalance(amount);
    }
}

```

 All methods (`deposit`, `withdraw`) execute within a transaction automatically.

## Method-Level `@Transactional`

- ✓ When applied at the **method level**, it affects **only that specific method**, overriding any class-level transactions.
- ✓ Useful for **customizing individual methods**, e.g., **different propagation types or rollback settings**.

### Example: Overriding `@Transactional` at Method Level

```
@Service
@Transactional // Default transaction applies to all methods
public class OrderService {

    public void placeOrder(Order order) {
        order.process();
    }

    @Transactional(propagation = Propagation.REQUIRES_NEW, rollbackFor = Exception.class)
    public void cancelOrder(Order order) throws Exception {
        order.reverseTransaction();
        throw new Exception("Order cancellation failed!");
    }
}
```

💡 The `cancelOrder()` method runs in a separate transaction and rolls back on all exceptions.

## Key Differences Between Class-Level & Method-Level `@Transactional`

Aspect	Class-Level <code>@Transactional</code>	Method-Level <code>@Transactional</code>
Scope	Applies to <b>all methods</b> in the class	Applies <b>only to the specific method</b>
Override	Can be overridden by method-level transaction	Overrides class-level configuration
Flexibility	Less customizable for individual methods	Allows <b>specific configurations</b> per method

## 9. How to handle distributed transaction management in spring boot micro services?

Handling ACID (Atomicity, Consistency, Isolation, Durability) in distributed transaction management is complex because it involves multiple systems or services potentially across different networks or databases. Here's how it's generally done:

1. Two-Phase Commit (2PC)
  - ✓ Ensures atomicity by coordinating transactions across multiple databases.
  - ✓ Uses a coordinator to prepare and commit transactions in two steps.
  - ✗ Downside: Can cause performance bottlenecks and blocking issues.
2. Saga Pattern
  - ✓ Breaks a transaction into multiple compensatable steps across microservices.
  - ✓ If a step fails, a compensating transaction rolls back previous changes.
  - ✗ Downside: Requires careful design to handle failures properly.
3. Eventual Consistency with Event-Driven Architecture
  - ✓ Uses asynchronous events to propagate changes across services.
  - ✓ Ensures consistency over time rather than immediate atomicity.
  - ✗ Downside: Requires monitoring and reconciliation mechanisms.

## 10. Explain distributed transaction management in micro services?

### Distributed Transaction Management in Microservices

In a **monolithic application**, managing a transaction is simple because all operations occur within the same database and JVM.

But in **microservices**, each service typically manages **its own database** — this makes **distributed transaction management** challenging.

---

### What is a Distributed Transaction?

A **distributed transaction** spans **multiple services and databases**, and all operations must either **commit or rollback** as a single unit.

### Challenges

- Each microservice is **independently deployed**, often using different databases.
- There's **no shared transaction context** across services.
- Network failures and partial failures are common.

## 2. Saga Pattern – Asynchronous, Eventual Consistency

Instead of a global transaction, a **Saga** is a sequence of **local transactions**, where each service:

- Performs a transaction
- Publishes an event
- The next service listens and acts
- If something fails, **compensating transactions** are triggered

```
Here's a simple Choreography Pattern example in Spring Boot using Kafka. The Choreography pattern is a decentralized communication style often used in microservices, where services react to events rather than being orchestrated by a central service.

Scenario: Order Processing System
We have 3 microservices:
Order Service - Creates an order and publishes an event.
Payment Service - Listens to the order-created event, processes payment, and publishes a payment-processed event.
Inventory Service - Listens to the payment-processed event and updates stock.

Technology Stack
Spring Boot
Apache Kafka
Kafka Topics
order-created
payment-processed

1. Order Service
Producer: Order Created Event
@Service
public class OrderService {
    @Autowired
    private KafkaTemplate<String, String> kafkaTemplate;
    public void createOrder(String orderId) {
        // Business logic to create order...
        kafkaTemplate.send("order-created", orderId);
    }
}
```

```

2. Payment Service
Consumer: Order Created Event

@Service
public class PaymentService {

    @Autowired
    private KafkaTemplate<String, String> kafkaTemplate;

    @KafkaListener(topics = "order-created", groupId = "payment-group")
    public void processPayment(String orderId) {
        // Business logic to process payment...
        kafkaTemplate.send("payment-processed", orderId);
    }
}

3. Inventory Service
Consumer: Payment Processed Event

@Service
public class InventoryService {

    @KafkaListener(topics = "payment-processed", groupId = "inventory-group")
    public void updateInventory(String orderId) {
        // Business logic to update inventory...
        System.out.println("Inventory updated for order: " + orderId);
    }
}

application.yml (for all services)

spring:
  kafka:
    bootstrap-servers: localhost:9092
    consumer:
      group-id: <appropriate-group-id>
      auto-offset-reset: earliest
      key-deserializer: org.apache.kafka.common.serialization.StringDeserializer
      value-deserializer: org.apache.kafka.common.serialization.StringDeserializer
    producer:
      key-serializer: org.apache.kafka.common.serialization.StringSerializer
      value-serializer: org.apache.kafka.common.serialization.StringSerializer

  Notes
  Each service is independent and communicates via Kafka topics.

  No centralized coordinator (orchestrator); services react to events.

  Use schemas (e.g., Avro, Protobuf, or JSON) for structured message exchange in production.

```

## 11. Explain ACID properties?

Property	Description
Atomicity	All or nothing — either the entire transaction succeeds, or none of it does
Consistency	Ensures data moves from one valid state to another
Isolation	Transactions don't interfere with each other
Durability	Once committed, the transaction is permanent

## 1 Atomicity (All or Nothing)

Ensures that a **transaction** is either fully completed or fully rolled back. If any step fails, **changes made so far are undone**.

### Example: Bank Money Transfer

```
import org.springframework.transaction.annotation.Transactional;
import org.springframework.stereotype.Service;

@Service
public class BankService {

    @Transactional
    public void transferMoney(Account from, Account to, double amount) {
        from.withdraw(amount);
        to.deposit(amount);

        if (amount > 10000) {
            throw new RuntimeException("Transaction exceeds limit!");
        }
    }
}
```

## 2 Consistency (Valid State Before & After Transactions)

Ensures that **database integrity constraints** are maintained before and after the transaction.

### Example: Enforcing Data Integrity

Java

 Copy

```
@Transactional
public void createOrder(Order order) {
    if (!order.hasValidCustomer()) {
        throw new RuntimeException("Invalid customer ID!");
    }
    order.save();
}
```

 Prevents saving orders that **violate integrity rules**, ensuring database consistency.

## 3 Isolation (Concurrency Control)

Defines **how transactions interact concurrently**, preventing data corruption.

### Example: Setting Transaction Isolation Level

Java

 Copy

```
@Transactional(isolation = Isolation.READ_COMMITTED)
public void checkBalance(Account account) {
    account.fetchBalance();
}
```

 Prevents **dirty reads** by ensuring changes from other transactions are visible only after commit.

## 4 Durability (Ensures Persistence)

Guarantees that **once a transaction is committed, changes are permanent**.

### Example: Writing to a Transaction Log

Java

 Copy

```
@Transactional
public void saveTransactionLog(TransactionLog log) {
    log.writeToDatabase();
}
```

 Even if the system crashes, committed transactions are **recoverable through transaction logs**.

## 12. What happens if a transaction times out in Spring?

 Spring **automatically rolls back** the transaction if it exceeds the configured timeout.

 Use `timeout` parameter in `@Transactional` to specify the max execution time.

### 💡 Example: Setting Transaction Timeout

Java

 Copy

```
@Transactional(timeout = 5) // Transaction times out after 5 seconds
public void longRunningProcess() {
    // Heavy database operations
}
```

 Ensures slow transactions do not block database resources.

### 13. What is the benefit of using `@Transactional(readOnly = true)`?

- Improves **performance** by avoiding locks on the database.
- Optimizes **read operations**, reducing transaction overhead.

 **Example:** Using `readOnly = true` for performance optimization

Java

 Copy

```
@Transactional(readOnly = true)
public List<Customer> fetchCustomers() {
    return customerRepository.findAll();
}
```

 Ideal for queries where no data modification occurs.

### 14. What happens if you call a transactional method inside another transactional method?

- If the called method uses `REQUIRED` propagation, it runs within the parent transaction.
- If the called method uses `REQUIRES_NEW` propagation, it starts a new transaction **independent of the parent**.
- If the parent transaction **rolls back**, transactions using `REQUIRED` also rollback, but transactions using `REQUIRES_NEW` remain committed.

 **Example:** Nested Transactions

Java

 Copy

```
@Transactional
public void transferMoney(Account from, Account to, double amount) {
    from.withdraw(amount);
    to.deposit(amount);

    logService.logTransaction(from, to, amount); // Uses REQUIRES_NEW
}
```

 If `transferMoney()` fails, balance changes roll back, but logging remains committed.

### 15. How does Spring handle rollback in transactions?

- ✓ By default, **unchecked exceptions** (`RuntimeException`, `Error`) **trigger rollback**.
- ✓ Checked exceptions (`Exception`) **do not trigger rollback**, unless explicitly set using `rollbackFor`.

 **Example: Rollback on Checked Exception**

Java

 Copy

```
@Transactional(rollbackFor = Exception.class)
public void placeOrder(Order order) throws Exception {
    order.processPayment();
    order.updateStock();
    if (!order.isValid()) {
        throw new Exception("Invalid Order!");
    }
}
```

 Without `rollbackFor = Exception.class`, **Spring will commit the transaction even if an exception occurs.**

## 16. What is the difference between Declarative and Programmatic Transaction Management?

Type	Declarative Transactions ( <code>@Transactional</code> )	Programmatic Transactions ( <code>PlatformTransactionManager</code> )
Implementation	Uses <code>@Transactional</code> annotation	Uses <code>TransactionTemplate</code> or <code>TransactionManager</code>
Control	Spring handles transactions automatically	Full control over transaction handling
Code Complexity	<b>Simpler</b> – No explicit commit/rollback code	<b>More complex</b> – Requires manual commit/rollback
Recommended For	Standard transactions (JPA, JDBC, Hibernate)	Advanced transaction handling (nested transactions, manual rollbacks)

 **Example: Programmatic Transaction Management**

Java

 Copy

```
TransactionTemplate transactionTemplate = new TransactionTemplate(transactionManager);
transactionTemplate.execute(status -> {
    try {
        account.withdraw(100);
        account.deposit(100);
        return true;
    } catch (Exception e) {
        status.setRollbackOnly(); // Manual rollback
        return false;
    }
});
```

## 17. Explain realtime Spring Boot Distributed Transactions with Kafka Example?

Distributed transactions ensure **data consistency** across multiple services/databases, especially in **microservices architecture**. Spring Boot supports **event-driven distributed transactions** using **Kafka** as the messaging broker.

### 📌 Key Concepts

- Two-Phase Commit (2PC)** – Ensures transaction consistency across multiple databases (Not recommended for performance).
- Saga Pattern** – Splits a transaction into **multiple compensating transactions** for eventual consistency.
- Event-Driven Transactions** – Uses **Kafka** to manage transactional events across services asynchronously.

### ◆ Example: Distributed Transaction in Order Processing

#### Scenario

1. **Order Service** creates an order and sends an event (`ORDER_CREATED`) to Kafka.
  2. **Payment Service** listens to the event, processes payment, and publishes `PAYMENT_COMPLETED`.
  3. **Inventory Service** listens to the event, reserves stock, and commits changes only if payment succeeds.
- If any step fails, a **compensating transaction** (rollback) occurs.

### Step 1: Configure Kafka Producer in Order Service

```
Java Copy
@Service
public class OrderService {
    private final KafkaTemplate<String, String> kafkaTemplate;

    public OrderService(KafkaTemplate<String, String> kafkaTemplate) {
        this.kafkaTemplate = kafkaTemplate;
    }

    public void createOrder(Order order) {
        // Save order in database
        System.out.println("Order Created: " + order.getId());

        // Send event to Kafka
        kafkaTemplate.send("order-topic", "ORDER_CREATED:" + order.getId())
    }
}
```

## Step 2: Consume Kafka Event in Payment Service

Java

 Copy

```
@KafkaListener(topics = "order-topic", groupId = "payment-group")
public void processPayment(String message) {
    System.out.println("Processing Payment for: " + message);

    // Simulating payment success
    boolean paymentSuccess = true;

    if (paymentSuccess) {
        System.out.println("Payment Completed: " + message);
        // Publish event for inventory update
    } else {
        System.out.println("Payment Failed! Rolling back ...");
    }
}
```

## Step 3: Inventory Service Handling Compensation

Java

 Copy

```
import org.springframework.stereotype.Service;

@Service
public class InventoryService {

    @KafkaListener(topics = "payment-topic", groupId = "inventory-group")
    public void reserveStock(String message) {
        System.out.println("Reserving Stock for: " + message);

        if (message.contains("PAYMENT_COMPLETED")) {
            System.out.println("Stock Reserved Successfully!");
        } else {
            System.out.println("Stock Reservation Failed! Rolling back ...");
        }
    }
}
```

### • Key Benefits of Kafka for Distributed Transactions

- Decouples Microservices** – Services operate independently through Kafka events.
- Supports Eventual Consistency** – Failsafe mechanisms allow retries and rollbacks.
- Improves Performance** – Eliminates blocking transactions across multiple databases.
- Enhances Scalability** – Easily handles high-load transaction processing.

## 16. Error handling strategies or event replay mechanisms in distributed transactions?

Distributed transactions require robust error handling and event replay mechanisms to ensure data consistency and fault tolerance in microservices. Here are key techniques:

### 1 Error Handling Strategies

Distributed systems often face **network failures, service downtime, and inconsistent states**. Effective error handling ensures reliability.

- Retry Mechanism** → Automatically retries failed transactions using **exponential backoff** to prevent overload.
- Dead Letter Queue (DLQ)** → Stores failed events in a **special Kafka queue** for later inspection and processing.
- Idempotency Keys** → Prevents duplicate transaction processing **by ensuring requests are uniquely identifiable**.
- Compensating Transactions (Rollback Actions)** → Undo operations via **Saga Pattern** when failures occur in multi-step transactions.

#### 💡 Example: Retrying Failed Transactions

Java

 Copy

```
@KafkaListener(topics = "order-topic", groupId = "payment-group")
public void processPayment(String message) {
    try {
        // Process payment
    } catch (Exception e) {
        retryTransaction(message); // Retry logic with backoff strategy
    }
}
```

### 2 Event Replay Mechanisms

Event replay allows **recovering lost or failed transactions by replaying Kafka events or reprocessing database logs**.

- Kafka Consumer Offsets** → Replay events from a specific offset using **consumer group checkpoints**.
- Event Sourcing** → Maintain an **event log** to reconstruct past transactions and restore system state.
- Outbox Pattern** → Stores events **in the database** and later publishes them reliably to Kafka for replay.
- Replay Dead Letter Events** → Periodically process **failed events** from DLQ and reinject them into the system.

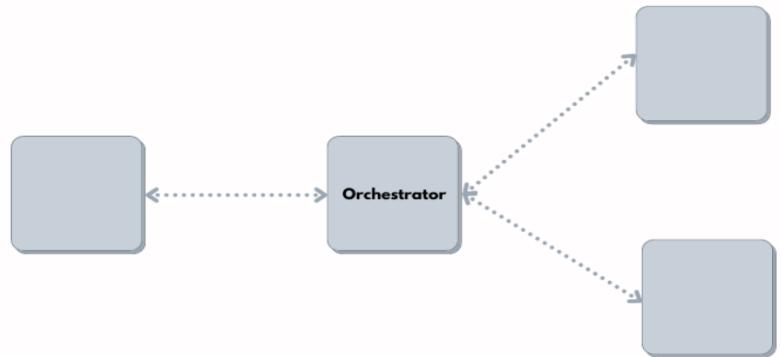
#### 💡 Example: Replaying Kafka Events

```
ConsumerRecords<String, String> records = kafkaConsumer.poll(Duration.ofSeconds(5));
for (ConsumerRecord<String, String> record : records) {
    processEvent(record.value()); // Replay stored transactions
}
```

## 18. Example for Orchestra pattern?

In this pattern, we will have an orchestrator, a separate service, which will be coordinating all the transactions among all the Microservices. If things are fine, it makes the order-request as complete, otherwise marks that as cancelled.

# Saga Pattern



## 🎯 Use Case: Order Workflow (Orchestrated)

CSS

```
Client → Order Orchestrator Service
    |--- calls Order Service
    |--- calls Inventory Service
    |--- calls Payment Service
```

If any step fails, the Orchestrator triggers **compensating actions**.

# E-commerce Architecture

## Microservices:

1. `order-service` - creates & cancels order
2. `inventory-service` - reserves & releases items
3. `payment-service` - charges & refunds
4. `orchestrator-service` - coordinates all steps

Each service exposes its own REST API and runs on a separate port.

## Example: Service Endpoints Overview

Service	Port	REST Endpoint
order-service	8081	<code>POST /orders</code> , <code>DELETE /orders/{id}</code>
inventory-service	8082	<code>POST /inventory/reserve</code> , <code>DELETE /inventory/release/{item}</code>
payment-service	8083	<code>POST /payment/charge</code> , <code>POST /payment/refund</code>
orchestrator	8080	<code>POST /orchestrator/orders</code>

## 💡 STEP-BY-STEP SETUP (abridged)

### ✓ order-service (port 8081)

```
java

@RestController
@RequestMapping("/orders")
public class OrderController {
    @PostMapping
    public ResponseEntity<String> createOrder(@RequestBody OrderRequest order) {
        return ResponseEntity.ok("Order created: " + order.getOrderId());
    }

    @DeleteMapping("/{id}")
    public ResponseEntity<String> cancelOrder(@PathVariable String id) {
        return ResponseEntity.ok("Order cancelled: " + id);
    }
}
```

.1.

### ✓ inventory-service (port 8082)

```
java

@RestController
@RequestMapping("/inventory")
public class InventoryController {
    @PostMapping("/reserve")
    public ResponseEntity<String> reserve(@RequestBody InventoryRequest req) {
        return ResponseEntity.ok("Inventory reserved: " + req.getItem());
    }

    @DeleteMapping("/release/{item}")
    public ResponseEntity<String> release(@PathVariable String item) {
        return ResponseEntity.ok("Inventory released: " + item);
    }
}
```

## payment-service (port 8083)

java

```
@RestController
@RequestMapping("/payment")
public class PaymentController {
    @PostMapping("/charge")
    public ResponseEntity<String> charge(@RequestBody PaymentRequest req) {
        if (req.getAmount() > 1000) throw new RuntimeException("Payment Failed");
        return ResponseEntity.ok("Payment charged: ₹" + req.getAmount());
    }

    @PostMapping("/refund")
    public ResponseEntity<String> refund(@RequestBody PaymentRequest req) {
        return ResponseEntity.ok("Payment refunded: ₹" + req.getAmount());
    }
}
```

## orchestrator-service (port 8080)

java

```
@Service
public class OrderOrchestratorService {

    private final RestTemplate restTemplate = new RestTemplate();

    public String placeOrder(OrderRequest order) {
        try {
            restTemplate.postForObject("http://localhost:8081/orders", order, String.class);
            restTemplate.postForObject("http://localhost:8082/inventory/reserve", order, String.class);
            restTemplate.postForObject("http://localhost:8083/payment/charge", order, String.class);
            return "✅ Order Completed";
        } catch (Exception e) {
            restTemplate.postForObject("http://localhost:8083/payment/refund", order, String.class);
            restTemplate.delete("http://localhost:8082/inventory/release/" + order.getItem());
            restTemplate.delete("http://localhost:8081/orders/" + order.getOrderId());
            return "❌ Failed, Compensated";
        }
    }
}
```

```
@RestController
@RequestMapping("/orchestrator/orders")
public class OrchestratorController {
    @Autowired private OrderOrchestratorService orchestrator;

    @PostMapping
    public ResponseEntity<String> create(@RequestBody OrderRequest order) {
        return ResponseEntity.ok(orchestrator.placeOrder(order));
    }
}
```

## Sample DTO

java

```
public class OrderRequest {
    private String orderId;
    private String item;
    private double amount;
    // Getters & Setters
}
```

## Testing via Curl/Postman

### Success:

json

```
POST http://localhost:8080/orchestrator/orders
{
    "orderId": "abc123",
    "item": "Book",
    "amount": 500
}
```

### Failure & rollback:

json

```
{
    "orderId": "abc124",
    "item": "Laptop",
    "amount": 1500
}
```

## Testing via Curl/Postman

### Success:

```
json

POST http://localhost:8080/orchestrator/orders
{
    "orderId": "abc123",
    "item": "Book",
    "amount": 500
}
```

### Failure & rollback:

```
json

{
    "orderId": "abc124",
    "item": "Laptop",
    "amount": 1500
}
```

## 19. Can we use kafka in orchestra pattern?

Yes, you can use Kafka with Orchestration, but it depends on the design goals.



## 1. When Kafka Is NOT Used in Orchestration

In most traditional **Saga Orchestration** patterns:

- The **orchestrator service** makes **direct REST calls** (synchronous).
- Services respond immediately.
- Rollback is handled directly if any step fails.

This is what we've already shown in the Spring Boot REST example.

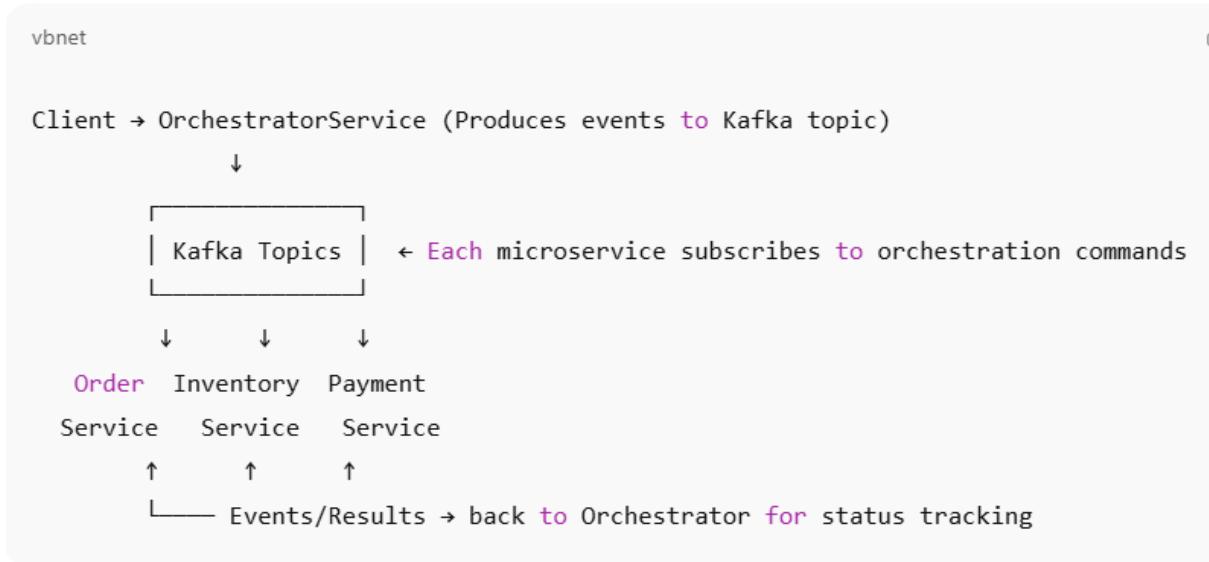


## 2. When Kafka IS Used in Orchestration

You might want to use Kafka-based orchestration when:

- You need **asynchronous execution** (non-blocking).
- You want **resilient and decoupled** communication between services.
- You want to **persist events**, enable **retries**, and improve **fault tolerance**.

### 🧠 Example Architecture Using Kafka for Orchestration:



- **Request Topics:** Orchestrator → Services
- **Response Topics:** Services → Orchestrator

### 20. Can we use rest implementation in choreography pattern?

Yes, you **can use REST in a Choreography pattern**, but it's **not recommended** for large-scale or complex workflows

## Example (REST-based Choreography)

Let's say:

1. **Order Service** creates an order  
→ it then calls **Inventory Service** via REST
2. **Inventory Service** reserves stock  
→ it then calls **Payment Service** via REST
3. If payment fails, **Inventory** and **Order** must compensate

This is choreography — no central orchestrator — but every service **initiates the next step**.

## Scenario: Simple Order Fulfillment (Choreography using REST)

Step	Microservice	Action
1	order-service	Creates the order and <b>calls inventory-service</b>
2	inventory-service	Reserves stock and <b>calls payment-service</b>
3	payment-service	Charges payment and <b>ends the flow</b>

All interactions are **peer-to-peer HTTP calls**, with **no central orchestrator**.

## DTO Used in All Services

```
java

public class OrderRequest {
    private String orderId;
    private String item;
    private double amount;
    // getters & setters
}
```

## 1 order-service (port 8081)

```
java

@RestController
@RequestMapping("/orders")
public class OrderController {

    private final RestTemplate restTemplate = new RestTemplate();

    @PostMapping
    public ResponseEntity<String> createOrder(@RequestBody OrderRequest order) {
        System.out.println("✓ Order created: " + order.getOrderId());

        // Call Inventory Service
        String inventoryUrl = "http://localhost:8082/inventory/reserve";
        restTemplate.postForEntity(inventoryUrl, order, String.class);

        return ResponseEntity.ok("Order processed");
    }
}
```

## 2 inventory-service (port 8082)

```
java

@RestController
@RequestMapping("/inventory")
public class InventoryController {

    private final RestTemplate restTemplate = new RestTemplate();

    @PostMapping("/reserve")
    public ResponseEntity<String> reserveInventory(@RequestBody OrderRequest order) {
        System.out.println("📦 Inventory reserved for item: " + order.getItem());

        // Call Payment Service
        String paymentUrl = "http://localhost:8083/payments/charge";
        restTemplate.postForEntity(paymentUrl, order, String.class);

        return ResponseEntity.ok("Inventory reserved");
    }
}
```

### 3 payment-service (port 8083)

java

```
@RestController
@RequestMapping("/payments")
public class PaymentController {

    @PostMapping("/charge")
    public ResponseEntity<String> charge(@RequestBody OrderRequest order) {
        System.out.println("💰 Payment charged: ₹" + order.getAmount());

        return ResponseEntity.ok("Payment successful");
    }
}
```

### Run and Test:

Send request to `order-service`:

bash

```
POST http://localhost:8081/orders
{
    "orderId": "ORD123",
    "item": "Laptop",
    "amount": 55000
}
```

Expected Logs:

bash

```
✓ Order created: ORD123
📦 Inventory reserved for item: Laptop
💰 Payment charged: ₹55000
```

## 21. How to handle the roll back scenario in 2PC, Orchestra and Choreography pattern?

Here's a **comprehensive breakdown** of how rollback is handled in:

- Two-Phase Commit (2PC)
- Saga Orchestration
- Saga Choreography

Each includes **code examples** with rollback logic using **Spring Boot**, and the differences in handling consistency.

## ⌚ 1. Two-Phase Commit (2PC) – Using MySQL + JTA

**Scenario:** `order-service` and `payment-service` must commit or rollback **together**.

### 2PC Setup

Add dependency (example: Atomikos or Narayana):

```
xml
Copy Edit

<!-- JTA 2PC with Atomikos -->
<dependency>
    <groupId>com.atomikos</groupId>
    <artifactId>transactions-jta</artifactId>
    <version>5.0.8</version>
</dependency>
```



## Spring Boot Example:

java

```
@Service
public class OrderService {

    @Autowired private OrderRepository orderRepo;
    @Autowired private PaymentRepository paymentRepo;

    @Transactional
    public void createOrderWithPayment(OrderRequest request) {
        orderRepo.save(new Order(...));

        // Simulate failure
        if (request.getAmount() > 1000) {
            throw new RuntimeException("Payment threshold exceeded");
        }

        paymentRepo.save(new Payment(...));
    }
}
```

## ⌚ Rollback:

If any exception occurs, both DB operations rollback.

2PC ensures **strong consistency** but suffers from **tight coupling** and **distributed locks**.



## 2. Saga Pattern – Orchestration Rollback

A **central orchestrator** tells each microservice what to do and what to undo on failure.

## Example: Spring Boot REST-based Orchestration with Compensation

### Step 1: Orchestrator Service

```
java Copy Edit
@Service
public class OrderOrchestrator {

    @Autowired private RestTemplate restTemplate;

    public void processOrder(OrderRequest request) {
        try {
            restTemplate.postForEntity("http://inventory/reserve", request, String.class);
            restTemplate.postForEntity("http://payment/charge", request, String.class);
        } catch (Exception ex) {
            // Compensation logic
            restTemplate.postForEntity("http://inventory/cancel", request, String.class);
            throw new RuntimeException("Saga failed and rollback triggered");
        }
    }
}
```



### Step 2: Inventory Cancel Endpoint

```
java
@PostMapping("/cancel")
public ResponseEntity<?> cancel(@RequestBody OrderRequest request) {
    // Undo inventory reservation
    return ResponseEntity.ok("Inventory rollback");
}
```

#### Rollback:

- Manual
- Controlled by orchestrator
- Easy to implement compensations



## 3. Saga Pattern – Choreography Rollback

Each service reacts to events. Rollbacks are **decentralized** and **local compensations** are triggered by failure events.

### Kafka-based Choreography Example (Simplified)

#### Inventory Service (listens to OrderCreated )

```
java Copy Edit  
  
@KafkaListener(topics = "OrderCreated")  
public void reserve(OrderRequest request) {  
    try {  
        // Reserve inventory  
        kafkaTemplate.send("InventoryReserved", request);  
    } catch (Exception e) {  
        kafkaTemplate.send("OrderFailed", request);  
    }  
}
```

#### Payment Service (listens to InventoryReserved )

```
java  
  
@KafkaListener(topics = "InventoryReserved")  
public void charge(OrderRequest request) {  
    try {  
        // Charge payment  
        kafkaTemplate.send("PaymentDone", request);  
    } catch (Exception e) {  
        // ⚙ Compensate Inventory  
        kafkaTemplate.send("InventoryCancel", request);  
        kafkaTemplate.send("OrderFailed", request);  
    }  
}
```

#### Inventory Listens for Rollback Events

```
java  
  
@KafkaListener(topics = "InventoryCancel")  
public void cancelInventory(OrderRequest request) {  
    // Release reserved stock  
}
```

